R and clinical/lab data

2026 edition

Matteo Miotto

2026-04-25

Contents

Welcome **5**

1 Dataframes, vectors and numbers **7**

1.1 Dataframes 7

1.2 Import file 7

1.3 Vectors 10

1.4 Save a dataframe 20

1.5 Home exercise 21

2 Strings, Factors and Dataframe Handling **23**

2.1 What we will see today 23

2.2 Strings and Factors 23

2.3 Dataframe Handling 35

2.4 Save your results 51

2.5 Home exercise 52

2.6 Bonus: how the clinical data were created 53

3 Data visualization **55**

3.1 Data visualization in R with ggplot2 55

3.2 Geom 56

3.3 Labs 69

3.4 Scales 70

3.5 Facet 76

3.6 Theme 78

3.7 Save the plots 80
3.8 Advanced examples 80

Welcome

Welcome to this course, here you can find the notebooks used in the PhD course of 2026.

Subscribe to receive notification when a new content is uploaded! NO SPAM, just updates.

Stay updated

Get notified when new material is published.

Powered by Buttondown

Chapter 1

Dataframes, vectors and numbers

1.1 Dataframes

The usual data you will deal with are tabular (rows x columns) data, usually improperly called "Excel files".

The standard formats for raw data files in data analyses are *.csv*, *.tsv* and *.txt* files; these should be our starting point.

1.2 Import file

The first thing to do is to import the file into R session. We can do it with several functions, but the standard one for *.csv* files is `read.csv`.

Let's see how to do so:

```
my_data <- read.csv("data/input_df.csv")
```

1.2.0.1 Notes about variables

Variables are the "memory" of an R session, they represent box where R stores data/values you want.

You can think them as boxes that have a name on them:

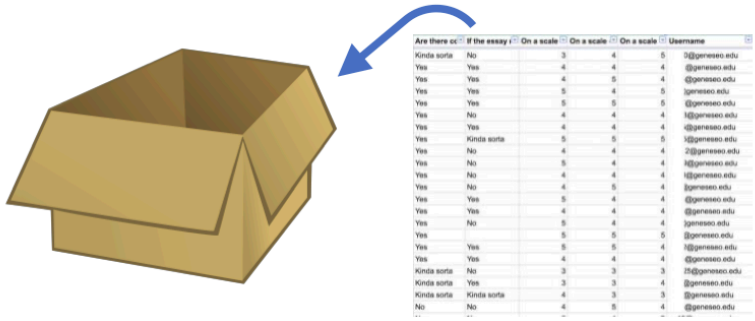

Are there co	If the essay i	On a scale	On a scale	On a scale	Username
Kinda sorta	No	3	4	5	3@geneseo.edu
Yes	Yes	4	4	4	@geneseo.edu
Yes	Yes	4	5	4	6@geneseo.edu
Yes	Yes	5	4	5	)geneseo.edu
Yes	Yes	5	5	5	6@geneseo.edu
Yes	No	4	4	4	1@geneseo.edu
Yes	Yes	4	4	4	4@geneseo.edu
Yes	Kinda sorta	5	5	5	1@geneseo.edu
Yes	No	4	4	4	2@geneseo.edu
Yes	No	5	4	4	8@geneseo.edu
Yes	No	4	4	4	1@geneseo.edu
Yes	No	4	5	4	)geneseo.edu
Yes	Yes	5	4	4	@geneseo.edu
Yes	Yes	4	4	4	@geneseo.edu
Yes	No	5	4	4	)geneseo.edu
Yes		5	5	5	)geneseo.edu
Yes	Yes	5	5	4	1@geneseo.edu
Yes	Yes	4	4	4	@geneseo.edu
Kinda sorta	No	3	3	3	75@geneseo.edu
Kinda sorta	Yes	3	3	4	)geneseo.edu
Kinda sorta	Kinda sorta	4	3	3	)geneseo.edu
No	No	4	5	4	@geneseo.edu

There are important **rules** about variables:

- In each variable you can store one single item (a number, a string, a dataframe, a list, ...)
- variable names can start only with a letter (uppercase or lowercase)
- variable names can contain **only** letters, numbers, underscores (_) and dots (.). **No spaces, no dash (-) or any other special characters**
- there could be only ONE variable with a specific name, saving an object with the same of another results in the substitution of the first one
- suggestion: variable names should be meaningful to you (and to the ones you share the code with): e.g. `x` is not a good name, while `patients_glucose_data` is preferred

1.2.1 First steps on a dataframe

When importing a dataframe, there is a fundamental step to do: *inspection*. Inspection helps us to understand how the dataframe is structured, which type

of data we have, which columns etc.

To do it, we use 3 functions that you need to remember (tatoo yourself): `summary`, `str` and `head`. Let's see them:

```
summary(my_data)
```

```
     sample        timepoint           genotype          treatment
 Min.   :  1.00   Length:192         Length:192         Length:192
 1st Qu.: 48.75   Class :character   Class :character   Class :character
 Median : 96.50   Mode  :character   Mode  :character   Mode  :character
 Mean   : 96.50
 3rd Qu.:144.25
 Max.   :192.00
      area        intensity
 Min.   :3000   Min.   : 9130286
 1st Qu.:3259   1st Qu.:16283904
 Median :3517   Median :18894103
 Mean   :3508   Mean   :19419175
 3rd Qu.:3769   3rd Qu.:22090201
 Max.   :3988   Max.   :29812114
```

```
str(my_data)
```

```
'data.frame':   192 obs. of  6 variables:
 $ sample   : int  1 2 3 4 5 6 7 8 9 10 ...
 $ timepoint: chr  "P1" "P1" "P1" "P1" ...
 $ genotype : chr  "wt" "wt" "wt" "wt" ...
 $ treatment: chr  "n" "n" "n" "n" ...
 $ area     : num  3125 3500 3773 3538 3936 ...
 $ intensity: num  16015465 19449815 21118210 18100075 20718554 ...
```

We can see how many rows and columns we have, as well as which columns, the data type of each and the first entries.

Lastly, `head` is the function that shows us the first 6 entries of the dataframe, as a table:

```
head(my_data)
```

```
  sample timepoint genotype treatment     area intensity
1      1        P1       wt         n 3124.709  16015465
2      2        P1       wt         n 3499.559  19449815
3      3        P1       wt         n 3773.000  21118210
4      4        P1       wt         n 3538.320  18100075
```

```
5       5        P1       wt         n 3936.014  20718554
6       6        P1       wt         n 3646.389  19507110
```

1.3 Vectors

Dataframes are organized and read by columns. Each column is a **vector** in R.

A vector is a collection of data of the **same type**; you can think of a vector as a pearl necklace, with each pearl being a single value:

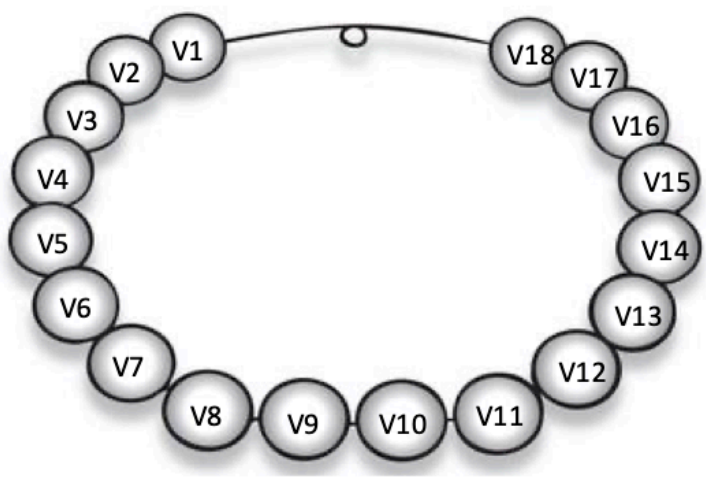

1.3.1 Accessing dataframe columns

For exercise purposes, we can create a new variable storing the values of the area column of the dataframe. To extract a single column from a dataframe, we use the structure `dataframe_name$column_name`.

```
area_values <- my_data$area
```

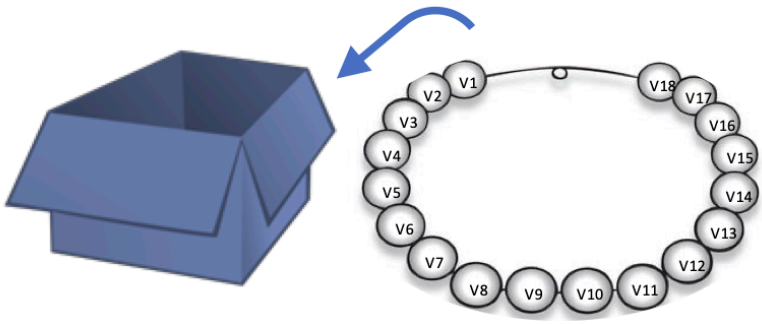

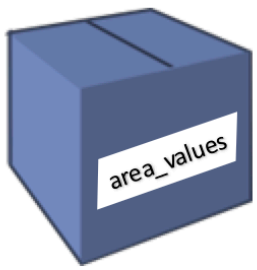

1.3.2 Accessing values of a vector: slicing

Slicing is the operation of extract values out of a variable (in this case a vector). It is done to answer these questions:

- *Which is the first value of the vector?*
- *Which is the value at position n of a vector?*
- *Which are the values from position n to m?*

There are different ways in which we can slice a vector, and today we will see indexing (through positions).

The syntax is:

- `variable_name[position]` to get one single value
- `variable_name[start:end]` to get elements from `start` to `end`

REMEMBER: In R, the first element is at index 1.

```
# extract one single value

first_area_value <- area_values[1]
print(first_area_value)
```

```
[1] 3124.709
```

```
# extract range of values
from_fifth_to_ninth_areas <- area_values[5:9]
print(from_fifth_to_ninth_areas)
```

```
[1] 3936.014 3646.389 3073.109 3717.063 3887.006
```

Q: *What 5:9 does in R?* A: It creates a vector containing numbers from 5 to 9, that's why it works:

```
print(5:9)
```

```
[1] 5 6 7 8 9
```

Q: *And to get last element?*

A: We need to know the lenght of the vector and use it as index. To get the length of a vector we use the function `length`:

```
length_area <- length(area_values)
print(length_area)
```

```
[1] 192
```

```
last_area <- area_values[length_area]
print(last_area)
```

```
[1] 3808.121
```

Exercise 1.1. Which is the 5th value of the intensity column of the dataframe?

Solution

We can do it in many steps:

```
intensity_values <- my_data$intensity

fifth_intensity <- intensity_values[5]

print(fifth_intensity)
```

```
[1] 20718554
```

Or in just one step:

```
fifth_intensity <- my_data$intensity[5]

print(fifth_intensity)
```

```
[1] 20718554
```

1.3.2.1 Replace a value by index

We can replace of a vector using indexes.

Note: new values should be of the same type (numbers for numbers, strings for strings ecc), otherwise everything will be changed to strings (most of the time).

To change a single value, we can type `variable_name[position] <- new_value`; to change multiple values `variable_name[start:end] <- new_values`.

1.3.3 Numbers in R

Now that we know how to get and extract numbers from numeric vectors, let's see what we can do with numbers in R.

1.3.3.1 Declare a new numeric variable

To declare a new numeric variable, we use the statement `variable_name <- number`:

```
conversion_factor <- 1000

print(conversion_factor)
```

```
[1] 1000
```

```r
print(typeof(conversion_factor))
```

```
[1] "double"
```

Decimal point numbers wants the dot `.` as decimal separator:

```r
decimal_n <- 215.6
print(decimal_n)
```

```
[1] 215.6
```

```r
print(typeof(decimal_n))
```

```
[1] "double"
```

1.3.3.2 Arithmetical operations

We can perform all sort of matematical operation:

```r
# sum

print(2 + 5)
```

```
[1] 7
```

```r
# subtraction

print(7 - 2.5)
```

```
[1] 4.5
```

```r
# multiplication
print(12 * 5)
```

```
[1] 60
```

```r
# division
print(105 / 43)
```

```
[1] 2.44186
```

```r
# exponential
print(5 ** 2)
```

```
[1] 25
```

```r
print(5 ^ 3)
```

```
[1] 125
```

```r
# square root
print(sqrt(12))
```

```
[1] 3.464102
```

We can combine them all to do more complex operations. For example, you can resolve this equation $x^2 - 7x + 12 = 0$. We know that the formula to resolve this equation is: $x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$, so we can reconstruct them in R:

```r
a <- 1
b <- -7
c <- 12

x1 <- (-b + sqrt(b**2 - 4*a*c)) / (2 * a)
x2 <- (-b - sqrt(b**2 - 4*a*c)) / (2 * a)

print(x1)
```

```
[1] 4
```

```r
print(x2)
```

```
[1] 3
```

Exercise 1.2. Which are the last 3 values of the intensity column of the dataframe?

Solution

We can do it in many steps:

```
intensity_values <- my_data$intensity

length_intensity <- length(intensity_values)

third_to_last_index <- length_intensity - 2

last_3_intensities <- intensity_values[third_to_last_index:length_intensity]

print(last_3_intensities)
```

```
[1] 22965783 22027811 22338574
```

1.3.4 Vector-number operations

What we have seen so far on arithmetical operations between single numbers, can be applied also to numerical vector x single number operations.

In our data, area is expressed in µm^2 and we want to convert it into mm^2. To do so, we should divide each value of the area by a factor of 1,000,000; the cool thing about R is that this is done automatically when we use the statement `numeric_vector <operand> number`:

```
division_factor = 1e6

areas_mm2 <- area_values / division_factor

print(head(areas_mm2))
```

```
[1] 0.003124709 0.003499559 0.003773000 0.003538320 0.003936014 0.003646389
```

```
print(head(area_values))
```

```
[1] 3124.709 3499.559 3773.000 3538.320 3936.014 3646.389
```

This is true for ANY arithmetical operations:

```
print(head(intensity_values))
```

```
[1] 16015465 19449815 21118210 18100075 20718554 19507110
```

```
print(head(intensity_values + 10))
```

```
[1] 16015475 19449825 21118220 18100085 20718564 19507120
```

1.3.5 Operations between vectors

It is possible also to perform element-wise operations between vectors.

Let's load a new dataset with some patient data, and calculate the BMI of each patient.

```r
patient_df <- read.csv("data/patient_data.csv")

str(patient_df)
```

```
'data.frame':   100 obs. of  3 variables:
 $ patient_id: int  1 2 3 4 5 6 7 8 9 10 ...
 $ weight    : num  75.8 61.8 91.7 66.4 80.4 ...
 $ height    : num  186 212 188 176 174 ...
```

```r
summary(patient_df)
```

```
   patient_id         weight           height      
 Min.   :  1.00   Min.   : 49.79   Min.   : 91.88  
 1st Qu.: 25.75   1st Qu.: 68.25   1st Qu.:159.20  
 Median : 50.50   Median : 73.91   Median :173.15  
 Mean   : 50.50   Mean   : 75.18   Mean   :174.96  
 3rd Qu.: 75.25   3rd Qu.: 84.61   3rd Qu.:185.90  
 Max.   :100.00   Max.   :112.18   Max.   :253.13  
```

```r
head(patient_df)
```

```
  patient_id   weight   height
1          1 75.82265 185.6147
2          2 61.77137 212.2587
3          3 91.65377 188.3604
4          4 66.38117 176.2986
5          5 80.42689 174.0509
6          6 79.32015 173.7906
```

BMI formula: $\frac{\text{weight in kg}}{(\text{height in m})^2}$

To do so, we should transform the height in m, and then perform the operation. We will save the results as a **new column** of the dataframe; this can be done with the syntax `dataframe_variable$new_column_name <- vector` or `dataframe_variable$["new_column_name"] <- vector` (we will see on day 2 the strings).

```r
# Get height in m
patient_df$height_in_m <- patient_df$height / 100

print(head(patient_df))
```

```
  patient_id   weight   height height_in_m
1          1 75.82265 185.6147    1.856147
2          2 61.77137 212.2587    2.122587
3          3 91.65377 188.3604    1.883604
4          4 66.38117 176.2986    1.762986
5          5 80.42689 174.0509    1.740509
6          6 79.32015 173.7906    1.737906
```

```r
patient_df["BMI"] <- patient_df$weight / (patient_df$height_in_m ** 2)

print(head(patient_df))
```

```
  patient_id   weight   height height_in_m      BMI
1          1 75.82265 185.6147    1.856147 22.00768
2          2 61.77137 212.2587    2.122587 13.71059
3          3 91.65377 188.3604    1.883604 25.83279
4          4 66.38117 176.2986    1.762986 21.35733
5          5 80.42689 174.0509    1.740509 26.54902
6          6 79.32015 173.7906    1.737906 26.26220
```

1.3.6 Sumamry statistics of numeric vectors

Usually, when dealing with numeric data we want to have some summary statistics on a specific data (e.g. mean values, median, quartile, standard deviation, …).

In R there are many built-in functions that can help us in doing so:

- `mean` to calculate the mean
- `sd` to calculate the standard deviation
- `variance` to calculate the variance
- `median` to calculate the median
- `quantile` to calculate the nth quantile of a distribution
- `round` to round decimal values to n decimal places
- `sum` to sum all the values of the vector

- many others

They all have in common the syntax: `name_of_the_function(vector)`.

```
intensity_mean <- mean(intensity_values)
print(intensity_mean)
```

```
[1] 19419175
```

```
height_sd <- sd(patient_df$height)
print(height_sd)
```

```
[1] 26.73183
```

```
BMI_median <- median(patient_df$BMI)
print(BMI_median)
```

```
[1] 25.23683
```

```
first_quartile_weight <- quantile(patient_df$weight, probs = 0.25)
print(first_quartile_weight)
```

```
     25%
68.25118
```

```
rounded_BMI <- round(patient_df$BMI, 1)
print(head(rounded_BMI))
```

```
[1] 22.0 13.7 25.8 21.4 26.5 26.3
```

Exercise 1.3. Scale BMI values (formula: $scaled_i = \frac{x_i - \overline{x}}{sd(x)}$)

Solution

Here is a step-by-step solution:

```
# calculate BMI mean
mean_BMI <- mean(patient_df$BMI)

# calculate BMI sd
sd_BMI <- sd(patient_df$BMI)

# Calculate scaled BMIs
patient_df$scaled_BMI <- (patient_df$BMI - mean_BMI) / sd_BMI

print(head(patient_df))
```

```
  patient_id   weight   height height_in_m      BMI  scaled_BMI
1          1 75.82265 185.6147    1.856147 22.00768 -0.44542706
2          2 61.77137 212.2587    2.122587 13.71059 -1.29224320
3          3 91.65377 188.3604    1.883604 25.83279 -0.05502959
4          4 66.38117 176.2986    1.762986 21.35733 -0.51180319
5          5 80.42689 174.0509    1.740509 26.54902  0.01807031
6          6 79.32015 173.7906    1.737906 26.26220 -0.01120341
```

That is exactly what the formula `scale` do underneath:

```r
# Calculate scaled BMIs
patient_df$scaled_BMI_function <- scale(patient_df$BMI)

print(head(patient_df))
```

```
  patient_id   weight   height height_in_m      BMI  scaled_BMI
1          1 75.82265 185.6147    1.856147 22.00768 -0.44542706
2          2 61.77137 212.2587    2.122587 13.71059 -1.29224320
3          3 91.65377 188.3604    1.883604 25.83279 -0.05502959
4          4 66.38117 176.2986    1.762986 21.35733 -0.51180319
5          5 80.42689 174.0509    1.740509 26.54902  0.01807031
6          6 79.32015 173.7906    1.737906 26.26220 -0.01120341
  scaled_BMI_function
1         -0.44542706
2         -1.29224320
3         -0.05502959
4         -0.51180319
5          0.01807031
6         -0.01120341
```

1.4 Save a dataframe

We are satisfied with this preliminary edit of the patients' data, so we can save the data to a dedicated file.

We use `write.csv` function:

```r
write.csv(patient_df, # what we want to save
          file = "output/patients_data_BMI.csv", # name of the file
          row.names = FALSE) # we'll see row names in day 2
```

1.5 Home exercise

For next time, if you want, you can try to do this exercise:

1. Starting from patient data (raw), load the file
2. Inspect it
3. You know that you have to give drug A to each patient so that the final concentration is 10 mg/kg (mg of drug every kg of patient). Calculate how much drug you should give to each patient.
4. Given that a single stock of drug A is 5 g, how many stocks you have to order?

Solution

```r
# 1. Load the file
patient_df <- read.csv("data/patient_data.csv")
```

```r
# 2. Inspect it
head(patient_df)
```

```
  patient_id   weight   height
1          1 75.82265 185.6147
2          2 61.77137 212.2587
3          3 91.65377 188.3604
4          4 66.38117 176.2986
5          5 80.42689 174.0509
6          6 79.32015 173.7906
```

```r
str(patient_df)
```

```
'data.frame':   100 obs. of  3 variables:
 $ patient_id: int  1 2 3 4 5 6 7 8 9 10 ...
 $ weight    : num  75.8 61.8 91.7 66.4 80.4 ...
 $ height    : num  186 212 188 176 174 ...
```

```r
summary(patient_df)
```

```
   patient_id         weight           height      
 Min.   :  1.00   Min.   : 49.79   Min.   : 91.88  
 1st Qu.: 25.75   1st Qu.: 68.25   1st Qu.:159.20  
 Median : 50.50   Median : 73.91   Median :173.15  
 Mean   : 50.50   Mean   : 75.18   Mean   :174.96  
 3rd Qu.: 75.25   3rd Qu.: 84.61   3rd Qu.:185.90  
 Max.   :100.00   Max.   :112.18   Max.   :253.13  
```

```r
# 3. Calculate drug for each patient

patient_df$drugA_quantity_mg <- patient_df$weight * 10
head(patient_df)
```

```
  patient_id   weight   height drugA_quantity_mg
1          1 75.82265 185.6147          758.2265
2          2 61.77137 212.2587          617.7137
3          3 91.65377 188.3604          916.5377
4          4 66.38117 176.2986          663.8117
5          5 80.42689 174.0509          804.2689
6          6 79.32015 173.7906          793.2015
```

```r
# 4. How many stocks you have to order?

total_drug_A_mg <- sum(patient_df$drugA_quantity_mg)

total_drug_A_g <- total_drug_A_mg / 1000

stock_weight <- 5

number_of_stocks <- total_drug_A_g/stock_weight

print(number_of_stocks)
```

```
[1] 15.03539
```

To be precise, we can use `ceiling` function that round to the upper integer:

```r
number_of_stocks <- ceiling(number_of_stocks)

print(number_of_stocks)
```

```
[1] 16
```

Chapter 2

Strings, Factors and Dataframe Handling

2.1 What we will see today

Today we will look at some patient data, and we will answer these questions:

- Are data consistent?
- Which is the mean and sd values of glucose in each treatment group? How many patients we have per treatment group?
- Which group has the highest mean creatinine? (2 different datasets)

2.2 Strings and Factors

Today we will work with a clinical dataset containing information about patients enrolled in a multi-site trial. Let's import it and, as always, **inspect it first**:

```r
library(tidyverse)
```

```
-- Attaching core tidyverse packages ------------------------ tidyverse 2.0.0 --
v dplyr     1.2.1     v readr     2.2.0
v forcats   1.0.1     v stringr   1.6.0
v ggplot2   4.0.3     v tibble    3.3.1
v lubridate 1.9.5     v tidyr     1.3.2
v purrr     1.2.2
```

```
-- Conflicts ------------------------------------------ tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to
```

```r
clinical_data <- read.csv("data/clinical_data.csv")
```

```r
str(clinical_data)
```

```
'data.frame':   240 obs. of  15 variables:
 $ patient_id      : chr  "PT-001" "PT-002" "PT-003" "PT-004" ...
 $ site            : chr  "Milan" "Rome" "Rome" "Turin " ...
 $ visit_date      : chr  "2023-07-20" "2020-08-15" "2023-10-15" "2020-03-06" ...
 $ age             : int  40 40 59 46 58 53 51 34 43 60 ...
 $ sex             : chr  "Male" "Male" "Female" "Male" ...
 $ smoking         : chr  "never" "Curent" "former" "Curent" ...
 $ treatment       : chr  "Drug_A" "Drug_B" "Drug_A" "Drug_A" ...
 $ stage           : chr  "II" "I" "I" "I" ...
 $ weight_kg       : num  85.9 100.9 65 95.9 75.7 ...
 $ height_cm       : num  173 185 166 182 168 ...
 $ sbp_mmhg        : int  164 143 152 142 160 150 137 146 151 159 ...
 $ dbp_mmhg        : int  101 102 98 88 96 107 96 106 94 104 ...
 $ glucose_mgdl    : int  113 105 110 120 131 78 108 93 106 110 ...
 $ cholesterol_mgdl: int  172 239 146 232 192 232 163 199 200 185 ...
 $ crp_ngml        : num  5.89 6.13 7.81 3.35 21.72 ...
```

```r
summary(clinical_data)
```

```
  patient_id            site            visit_date             age
 Length:240         Length:240         Length:240         Min.   :26.0
 Class :character   Class :character   Class :character   1st Qu.:49.0
 Mode  :character   Mode  :character   Mode  :character   Median :57.5
                                                          Mean   :57.3
                                                          3rd Qu.:65.0
                                                          Max.   :85.0
     sex              smoking           treatment           stage
 Length:240         Length:240         Length:240         Length:240
 Class :character   Class :character   Class :character   Class :character
 Mode  :character   Mode  :character   Mode  :character   Mode  :character



   weight_kg        height_cm        sbp_mmhg        dbp_mmhg
 Min.   : 39.40   Min.   :148.9   Min.   :115.0   Min.   : 73.00
```

```
 1st Qu.: 65.35   1st Qu.:162.2   1st Qu.:139.0   1st Qu.: 90.75
 Median : 75.65   Median :168.5   Median :146.5   Median : 97.00
 Mean   : 76.87   Mean   :169.2   Mean   :146.7   Mean   : 96.11
 3rd Qu.: 87.83   3rd Qu.:174.8   3rd Qu.:153.2   3rd Qu.:102.00
 Max.   :132.90   Max.   :200.1   Max.   :192.0   Max.   :120.00
  glucose_mgdl   cholesterol_mgdl    crp_ngml
 Min.   : 60.0   Min.   :113.0    Min.   : 1.92
 1st Qu.: 98.0   1st Qu.:189.8    1st Qu.: 5.30
 Median :109.0   Median :207.0    Median : 8.37
 Mean   :108.8   Mean   :207.9    Mean   :10.35
 3rd Qu.:118.0   3rd Qu.:226.0    3rd Qu.:13.02
 Max.   :149.0   Max.   :307.0    Max.   :42.09
```

```r
head(clinical_data)
```

```
  patient_id   site visit_date age    sex smoking treatment stage weight_kg
1     PT-001  Milan 2023-07-20  40   Male   never    Drug_A    II      85.9
2     PT-002   Rome 2020-08-15  40   Male  Curent    Drug_B     I     100.9
3     PT-003   Rome 2023-10-15  59 Female  former    Drug_A     I      65.0
4     PT-004 Turin  2020-03-06  46   Male  Curent    Drug_A     I      95.9
5     PT-005  Turin 2023-03-30  58   Male  FORMER    Drug_B   III      75.7
6     PT-006  Milan 2023-04-15  53 Female  Curent   Control    IV      69.6
  height_cm sbp_mmhg dbp_mmhg glucose_mgdl cholesterol_mgdl crp_ngml
1     172.7      164      101          113              172     5.89
2     185.3      143      102          105              239     6.13
3     165.6      152       98          110              146     7.81
4     181.5      142       88          120              232     3.35
5     168.1      160       96          131              192    21.72
6     174.3      150      107           78              232    32.42
```

You will notice that several columns contain text: `patient_id`, `sex`, `treatment`, `stage`, `smoking`, `site`, `diagnosis_code`. In R, text data can be represented in two fundamentally different ways: as **strings** (also called *characters*) or as **factors**. Let's understand the difference.

A **string** (or character) is simply a sequence of text. R treats it as a raw piece of text and does not impose any structure on it.

A **factor** is a *categorical variable*: it has a fixed, predefined set of allowed values called **levels**. Under the hood, R stores factors as integers (one integer per level) and keeps a table that maps integers to labels. This makes factors memory-efficient and very useful for statistical modelling (e.g., in regression,

factors are automatically handled as categorical predictors) and for plotting purposes.

When to use strings and when factors?

Use **string** when…	Use **factor** when…
Text is free-form or unpredictable	The variable takes a limited set of known values
You need to manipulate the text	Order matters (e.g., stage I < II < III < IV)
Unique identifiers (patient IDs)	Input to statistical models

As some of our columns fall in the conditions for factors, let's convert them with the function `factor`:

```
clinical_data$sex       <- factor(clinical_data$sex)
clinical_data$treatment <- factor(clinical_data$treatment)
clinical_data$site      <- factor(clinical_data$site)
clinical_data$smoking   <- factor(clinical_data$smoking)

# For stage, the categories have a natural ORDER: we declare an ordered factor
clinical_data$stage <- factor(clinical_data$stage,
                              levels = c("I", "II", "III", "IV"),
                              ordered = TRUE)

str(clinical_data)
```

```
'data.frame':   240 obs. of  15 variables:
 $ patient_id      : chr  "PT-001" "PT-002" "PT-003" "PT-004" ...
 $ site            : Factor w/ 4 levels "Milan","Rome",..: 1 2 2 4 3 1 2 2 1 2 ...
 $ visit_date      : chr  "2023-07-20" "2020-08-15" "2023-10-15" "2020-03-06" ...
 $ age             : int  40 40 59 46 58 53 51 34 43 60 ...
 $ sex             : Factor w/ 2 levels "Female","Male": 2 2 1 2 2 1 2 1 2 2 ...
 $ smoking         : Factor w/ 7 levels "Curent","current",..: 5 1 3 1 4 1 3 6 5 4 ...
 $ treatment       : Factor w/ 3 levels "Control","Drug_A",..: 2 3 2 2 3 1 2 1 3 2 ...
 $ stage           : Ord.factor w/ 4 levels "I"<"II"<"III"<..: 2 1 1 1 3 4 1 1 3 2 ...
 $ weight_kg       : num  85.9 100.9 65 95.9 75.7 ...
 $ height_cm       : num  173 185 166 182 168 ...
 $ sbp_mmhg        : int  164 143 152 142 160 150 137 146 151 159 ...
 $ dbp_mmhg        : int  101 102 98 88 96 107 96 106 94 104 ...
 $ glucose_mgdl    : int  113 105 110 120 131 78 108 93 106 110 ...
 $ cholesterol_mgdl: int  172 239 146 232 192 232 163 199 200 185 ...
 $ crp_ngml        : num  5.89 6.13 7.81 3.35 21.72 ...
```

Now `summary` gives us much more useful information on those columns:

```r
summary(clinical_data)
```

```
 patient_id            site     visit_date             age           sex
Length:240         Milan :96   Length:240         Min.   :26.0   Female:118
Class :character   Rome  :77   Class :character   1st Qu.:49.0   Male  :122
Mode  :character   Turin :31   Mode  :character   Median :57.5
                   Turin :36                      Mean   :57.3
                                                  3rd Qu.:65.0
                                                  Max.   :85.0

        smoking     treatment  stage      weight_kg        height_cm
Curent      :35   Control:74   I  :76   Min.   : 39.40   Min.   :148.9
current     :28   Drug_A :88   II :73   1st Qu.: 65.35   1st Qu.:162.2
former      :44   Drug_B :78   III:53   Median : 75.65   Median :168.5
FORMER      :46                IV :38   Mean   : 76.87   Mean   :169.2
never       :26                         3rd Qu.: 87.83   3rd Qu.:174.8
Never       :28                         Max.   :132.90   Max.   :200.1
never smoker:33
   sbp_mmhg        dbp_mmhg       glucose_mgdl   cholesterol_mgdl
Min.   :115.0   Min.   : 73.00   Min.   : 60.0   Min.   :113.0
1st Qu.:139.0   1st Qu.: 90.75   1st Qu.: 98.0   1st Qu.:189.8
Median :146.5   Median : 97.00   Median :109.0   Median :207.0
Mean   :146.7   Mean   : 96.11   Mean   :108.8   Mean   :207.9
3rd Qu.:153.2   3rd Qu.:102.00   3rd Qu.:118.0   3rd Qu.:226.0
Max.   :192.0   Max.   :120.00   Max.   :149.0   Max.   :307.0

   crp_ngml
Min.   : 1.92
1st Qu.: 5.30
Median : 8.37
Mean   :10.35
3rd Qu.:13.02
Max.   :42.09
```

Notice how `summary` now shows the count per level for factors (e.g. how many Male vs Female), while before it just said `character`.

Moreover, we can see that in smoking column, we have different labels with the same meaning (never, Never, never smoker for never). R (as well as any programming language) is case sensitive, so it treats those 3 as different labels, even if for us they mean the same thing.

IMPORTANT: labels should be consistent across datasets when you fill them.

As we want to edit those to be consistent, we need to change the column back to be a character, with `as.character` function:

```
clinical_data$smoking   <- as.character(clinical_data$smoking)
str(clinical_data)
```

```
'data.frame':   240 obs. of  15 variables:
 $ patient_id      : chr  "PT-001" "PT-002" "PT-003" "PT-004" ...
 $ site            : Factor w/ 4 levels "Milan","Rome",..: 1 2 2 4 3 1 2 2 1 2 ...
 $ visit_date      : chr  "2023-07-20" "2020-08-15" "2023-10-15" "2020-03-06" ...
 $ age             : int  40 40 59 46 58 53 51 34 43 60 ...
 $ sex             : Factor w/ 2 levels "Female","Male": 2 2 1 2 2 1 2 1 2 2 ...
 $ smoking         : chr  "never" "Curent" "former" "Curent" ...
 $ treatment       : Factor w/ 3 levels "Control","Drug_A",..: 2 3 2 2 3 1 2 1 3 2 ...
 $ stage           : Ord.factor w/ 4 levels "I"<"II"<"III"<..: 2 1 1 1 3 4 1 1 3 2 ...
 $ weight_kg       : num  85.9 100.9 65 95.9 75.7 ...
 $ height_cm       : num  173 185 166 182 168 ...
 $ sbp_mmhg        : int  164 143 152 142 160 150 137 146 151 159 ...
 $ dbp_mmhg        : int  101 102 98 88 96 107 96 106 94 104 ...
 $ glucose_mgdl    : int  113 105 110 120 131 78 108 93 106 110 ...
 $ cholesterol_mgdl: int  172 239 146 232 192 232 163 199 200 185 ...
 $ crp_ngml        : num  5.89 6.13 7.81 3.35 21.72 ...
```

We can now proceed.

2.2.1 Operations on strings

The main package for string operations in R (and the one we will use throughout the course) is `stringr`, which is part of the `tidyverse`. All its functions start with `str_` to make them easy to remember.

2.2.1.1 Character substitution

Let's start by looking at how to change characters in a string: from "Curent" to "current" using `str_replace(vector, old_pattern, new_pattern)`.

```
print(unique(clinical_data$smoking))
```

```
[1] "never"        "Curent"       "former"       "FORMER"       "Never"
[6] "never smoker" "current"
```

```
clinical_data$smoking <- str_replace(clinical_data$smoking, "Curent", "current")
print(unique(clinical_data$smoking))
```

```
[1] "never"        "current"      "former"       "FORMER"       "Never"
[6] "never smoker"
```

Then we can unify "FORMER" and "former" by transforming the former to lowercase with `str_to_lower`:

```
print(unique(clinical_data$smoking))
```

```
[1] "never"        "current"      "former"       "FORMER"       "Never"
[6] "never smoker"
```

```
clinical_data$smoking <- str_to_lower(clinical_data$smoking)
print(unique(clinical_data$smoking))
```

```
[1] "never"        "current"      "former"       "never smoker"
```

We have converted both FORMER to former and Never to never. We lastly have to change "never smoker" to "never".

Exercise 2.1. Change "never smoker" to "smoker" and make smoker as factor.

Solution

```
clinical_data$smoking <- str_replace(clinical_data$smoking, "never smoker", "never")
clinical_data$smoking <- factor(clinical_data$smoking, levels = c("never", "former", "current"))

summary(clinical_data)
```

```
  patient_id            site     visit_date             age           sex
 Length:240         Milan :96   Length:240         Min.   :26.0   Female:118
 Class :character   Rome  :77   Class :character   1st Qu.:49.0   Male  :122
 Mode  :character   Turin :31   Mode  :character   Median :57.5
                    Turin :36                      Mean   :57.3
                                                   3rd Qu.:65.0
                                                   Max.   :85.0
    smoking     treatment  stage      weight_kg        height_cm
 never  :87   Control:74   I  :76   Min.   : 39.40   Min.   :148.9
 former :90   Drug_A :88   II :73   1st Qu.: 65.35   1st Qu.:162.2
 current:63   Drug_B :78   III:53   Median : 75.65   Median :168.5
                           IV :38   Mean   : 76.87   Mean   :169.2
                                    3rd Qu.: 87.83   3rd Qu.:174.8
                                    Max.   :132.90   Max.   :200.1
    sbp_mmhg        dbp_mmhg       glucose_mgdl   cholesterol_mgdl
```

```
 Min.   :115.0   Min.   : 73.00   Min.   : 60.0   Min.   :113.0
 1st Qu.:139.0   1st Qu.: 90.75   1st Qu.: 98.0   1st Qu.:189.8
 Median :146.5   Median : 97.00   Median :109.0   Median :207.0
 Mean   :146.7   Mean   : 96.11   Mean   :108.8   Mean   :207.9
 3rd Qu.:153.2   3rd Qu.:102.00   3rd Qu.:118.0   3rd Qu.:226.0
 Max.   :192.0   Max.   :120.00   Max.   :149.0   Max.   :307.0
    crp_ngml
 Min.   : 1.92
 1st Qu.: 5.30
 Median : 8.37
 Mean   :10.35
 3rd Qu.:13.02
 Max.   :42.09
```

For sake of completeness, we could have done all of this with a single function `replace_values` (for new versions of R):

```
clinical_data$smoking <- replace_values(as.character(clinical_data$smoking), c("never s
                                        c("Curent") ~ "current",
                                        c("FORMER") ~ "former")

clinical_data$smoking = factor(clinical_data$smoking)
```

If it doesn't work, you could use:

```
clinical_data = mutate(smoking = case_when(
    smoking %in% c("never smoker", "Never") ~ "never",
    smoking %in% c("FORMER", "Former", "Fomer") ~ "former",
    smoking == "Curent" ~ "current",
    .default = smoking),
    .data = clinical_data
    )

clinical_data$smoking = factor(clinical_data$smoking)
```

Q: *Why there are 2 Turin?* A: Let's check in details the possible labels of that column with function `levels`:

```
print(levels(clinical_data$site))
```

```
[1] "Milan"  "Rome"   "Turin"  "Turin "
```

Here's the hitch, some have a trailing space after Turin. We can get rid of them using `str_trim` function:

```r
print(levels(clinical_data$site))
```

```
[1] "Milan"  "Rome"   "Turin"  "Turin "
```

```r
clinical_data$site <- factor(str_trim(clinical_data$site))

print(levels(clinical_data$site))
```

```
[1] "Milan" "Rome"  "Turin"
```

```r
print(summary(clinical_data))
```

```
  patient_id           site     visit_date             age           sex      
 Length:240         Milan:96   Length:240         Min.   :26.0   Female:118  
 Class :character   Rome :77   Class :character   1st Qu.:49.0   Male  :122  
 Mode  :character   Turin:67   Mode  :character   Median :57.5               
                                                   Mean   :57.3               
                                                   3rd Qu.:65.0               
                                                   Max.   :85.0               
    smoking     treatment  stage      weight_kg        height_cm     
 current:63   Control:74   I  :76   Min.   : 39.40   Min.   :148.9  
 former :90   Drug_A :88   II :73   1st Qu.: 65.35   1st Qu.:162.2  
 never  :87   Drug_B :78   III:53   Median : 75.65   Median :168.5  
                           IV :38   Mean   : 76.87   Mean   :169.2  
                                    3rd Qu.: 87.83   3rd Qu.:174.8  
                                    Max.   :132.90   Max.   :200.1  
    sbp_mmhg        dbp_mmhg       glucose_mgdl   cholesterol_mgdl
 Min.   :115.0   Min.   : 73.00   Min.   : 60.0   Min.   :113.0   
 1st Qu.:139.0   1st Qu.: 90.75   1st Qu.: 98.0   1st Qu.:189.8   
 Median :146.5   Median : 97.00   Median :109.0   Median :207.0   
 Mean   :146.7   Mean   : 96.11   Mean   :108.8   Mean   :207.9   
 3rd Qu.:153.2   3rd Qu.:102.00   3rd Qu.:118.0   3rd Qu.:226.0   
 Max.   :192.0   Max.   :120.00   Max.   :149.0   Max.   :307.0   
    crp_ngml      
 Min.   : 1.92  
 1st Qu.: 5.30  
 Median : 8.37  
 Mean   :10.35  
 3rd Qu.:13.02  
 Max.   :42.09  
```

2.2.1.2 Other useful functions to handle strings

A few more `stringr` functions you will find very useful:

Function	What it does
`str_detect(x, pattern)`	Returns TRUE/FALSE: does the string contain the pattern?
`str_starts(x, pattern)`	Does the string *start* with the pattern?
`str_ends(x, pattern)`	Does the string *end* with the pattern?
`str_pad(x, width, side, pad)`	Pads a string to a given width
`str_sub(x, start, end)`	Extracts a substring by position

2.2.2 Operations on multiple columns

2.2.2.1 Split a column and merge multiple columns

Sometimes a single column contains multiple pieces of information that could be split into separate columns. Look at `visit_date`:

```
head(clinical_data$visit_date)
```

```
[1] "2023-07-20" "2020-08-15" "2023-10-15" "2020-03-06" "2023-03-30"
[6] "2023-04-15"
```

We can see that the year of visit varies a lot, maybe we should take it into account. We can split this into 3 separate columns using `separate` function:

```
clinical_data <- separate(clinical_data,
                          col = visit_date,         # column to split
                          into = c("visit_year", "visit_month", "visit_day"), # new c
                          sep = "-",                # separator
                          remove = FALSE)           # keep the original column

head(clinical_data)
```

```
  patient_id  site visit_date visit_year visit_month visit_day age    sex
1     PT-001 Milan 2023-07-20       2023          07        20  40   Male
2     PT-002  Rome 2020-08-15       2020          08        15  40   Male
3     PT-003  Rome 2023-10-15       2023          10        15  59 Female
4     PT-004 Turin 2020-03-06       2020          03        06  46   Male
5     PT-005 Turin 2023-03-30       2023          03        30  58   Male
6     PT-006 Milan 2023-04-15       2023          04        15  53 Female
  smoking treatment stage weight_kg height_cm sbp_mmhg dbp_mmhg glucose_mgdl
```

```
1   never    Drug_A    II      85.9     172.7      164      101          113
2 current    Drug_B     I     100.9     185.3      143      102          105
3  former    Drug_A     I      65.0     165.6      152       98          110
4 current    Drug_A     I      95.9     181.5      142       88          120
5  former    Drug_B   III      75.7     168.1      160       96          131
6 current   Control    IV      69.6     174.3      150      107           78
  cholesterol_mgdl crp_ngml
1              172     5.89
2              239     6.13
3              146     7.81
4              232     3.35
5              192    21.72
6              232    32.42
```

The opposite operation — merging multiple columns into one — is done with `unite`:

```r
# example: unite back together
clinical_data <- unite(clinical_data,
                       col = "visit_year_month",         # name of the new column
                       visit_year, visit_month,           # columns to merge
                       sep = "-",                       # separator
                       remove = FALSE)                  # keep original columns

head(clinical_data)
```

```
  patient_id  site visit_date visit_year_month visit_year visit_month visit_day
1     PT-001 Milan 2023-07-20          2023-07       2023          07        20
2     PT-002  Rome 2020-08-15          2020-08       2020          08        15
3     PT-003  Rome 2023-10-15          2023-10       2023          10        15
4     PT-004 Turin 2020-03-06          2020-03       2020          03        06
5     PT-005 Turin 2023-03-30          2023-03       2023          03        30
6     PT-006 Milan 2023-04-15          2023-04       2023          04        15
  age    sex smoking treatment stage weight_kg height_cm sbp_mmhg dbp_mmhg
1  40   Male   never    Drug_A    II      85.9     172.7      164      101
2  40   Male current    Drug_B     I     100.9     185.3      143      102
3  59 Female  former    Drug_A     I      65.0     165.6      152       98
4  46   Male current    Drug_A     I      95.9     181.5      142       88
5  58   Male  former    Drug_B   III      75.7     168.1      160       96
6  53 Female current   Control    IV      69.6     174.3      150      107
  glucose_mgdl cholesterol_mgdl crp_ngml
1          113              172     5.89
2          105              239     6.13
3          110              146     7.81
4          120              232     3.35
```

```
5          131              192    21.72
6           78              232    32.42
```

2.2.3 Column and row names

Every dataframe has **column names** (the variable names) and **row names** (by default a number, but we can change it).

To retrieve column names:

```r
colnames(clinical_data)
```

```
 [1] "patient_id"       "site"             "visit_date"       "visit_year_month"
 [5] "visit_year"       "visit_month"      "visit_day"        "age"
 [9] "sex"              "smoking"          "treatment"        "stage"
[13] "weight_kg"        "height_cm"        "sbp_mmhg"         "dbp_mmhg"
[17] "glucose_mgdl"     "cholesterol_mgdl" "crp_ngml"
```

To rename specific columns, we can assign a new vector of names to `colnames`, or use the tidyverse function `rename` (we will see this with `dplyr` below):

```r
clinical_data = rename(clinical_data, "diastolic_bp_mmhg" = "dbp_mmhg", "systolic_bp_mm

colnames(clinical_data)
```

```
 [1] "patient_id"        "site"              "visit_date"
 [4] "visit_year_month"  "visit_year"        "visit_month"
 [7] "visit_day"         "age"               "sex"
[10] "smoking"           "treatment"         "stage"
[13] "weight_kg"         "height_cm"         "systolic_bp_mmhg"
[16] "diastolic_bp_mmhg" "glucose_mgdl"      "cholesterol_mgdl"
[19] "crp_ngml"
```

For row names:

```r
# by default, row names are row numbers
head(rownames(clinical_data))
```

```
[1] "1" "2" "3" "4" "5" "6"
```

```
# we could set patient_id as row names
rownames(clinical_data) <- clinical_data$patient_id

head(rownames(clinical_data))
```

```
[1] "PT-001" "PT-002" "PT-003" "PT-004" "PT-005" "PT-006"
```

However, **be careful** with row names: they are not a proper column of the dataframe, so they don't follow along easily in many tidyverse operations. It is generally better to keep your identifiers as a regular column.

2.3 Dataframe Handling

The main verbs we will learn today are:

- `mutate` — add or modify columns
- `select` — choose which columns to keep
- `filter`/`filter_out` — keep/remove rows that satisfy a condition
- `group_by` + `summarize` / `reframe` — aggregate data by groups
- `left_join` / `inner_join` — merge two dataframes

2.3.1 Create new columns with mutate

`mutate` creates new columns or modifies existing ones.

A key feature of `dplyr` is the **pipe operator** `|>` (or `%>%` in older tidyverse style). It passes the result of one operation directly to the next, making code much more readable:

You can read `|>` as *"and then"*.

```
# Add a BMI column (we already know this formula!)
clinical_data <- clinical_data |> # take clinical_data and then
  mutate(height_m = height_cm / 100,
         BMI = weight_kg / height_m^2,
         BMI = round(BMI, 1))

head(clinical_data)
```

```
       patient_id  site visit_date visit_year_month visit_year visit_month
PT-001     PT-001 Milan 2023-07-20          2023-07       2023          07
PT-002     PT-002  Rome 2020-08-15          2020-08       2020          08
PT-003     PT-003  Rome 2023-10-15          2023-10       2023          10
PT-004     PT-004 Turin 2020-03-06          2020-03       2020          03
PT-005     PT-005 Turin 2023-03-30          2023-03       2023          03
PT-006     PT-006 Milan 2023-04-15          2023-04       2023          04
       visit_day age    sex smoking treatment stage weight_kg height_cm
PT-001        20  40   Male   never    Drug_A    II      85.9     172.7
PT-002        15  40   Male current    Drug_B     I     100.9     185.3
PT-003        15  59 Female  former    Drug_A     I      65.0     165.6
PT-004        06  46   Male current    Drug_A     I      95.9     181.5
PT-005        30  58   Male  former    Drug_B   III      75.7     168.1
PT-006        15  53 Female current   Control    IV      69.6     174.3
       systolic_bp_mmhg diastolic_bp_mmhg glucose_mgdl cholesterol_mgdl
PT-001              164               101          113              172
PT-002              143               102          105              239
PT-003              152                98          110              146
PT-004              142                88          120              232
PT-005              160                96          131              192
PT-006              150               107           78              232
       crp_ngml height_m  BMI
PT-001     5.89    1.727 28.8
PT-002     6.13    1.853 29.4
PT-003     7.81    1.656 23.7
PT-004     3.35    1.815 29.1
PT-005    21.72    1.681 26.8
PT-006    32.42    1.743 22.9
```

`mutate` can also create columns based on conditions, using `if_else` or `case_when`:

```r
# Classify BMI into categories
clinical_data <- clinical_data |>
  mutate(BMI_class = case_when(
    BMI < 18.5 ~ "Underweight",
    BMI < 25   ~ "Normal",
    BMI < 30   ~ "Overweight",
    BMI >= 30  ~ "Obese"
  ),
  BMI_class = factor(BMI_class, labels = c("Underweight", "Normal", "Overweight", "Obes
  )

summary(clinical_data$BMI_class)
```

```
Underweight      Normal  Overweight       Obese
         67          58         102          13
```

2.3.2 select

`select` chooses which **columns** to keep (or drop). This is useful when working with wide dataframes where you only need a subset of variables.

```r
# keep only demographic and clinical variables
demo_cols <- clinical_data |>
  select(patient_id, site, age, sex, treatment, stage, BMI, BMI_class)

head(demo_cols)
```

```
       patient_id  site age    sex treatment stage  BMI   BMI_class
PT-001     PT-001 Milan  40   Male    Drug_A    II 28.8  Overweight
PT-002     PT-002  Rome  40   Male    Drug_B     I 29.4  Overweight
PT-003     PT-003  Rome  59 Female    Drug_A     I 23.7 Underweight
PT-004     PT-004 Turin  46   Male    Drug_A     I 29.1  Overweight
PT-005     PT-005 Turin  58   Male    Drug_B   III 26.8  Overweight
PT-006     PT-006 Milan  53 Female   Control    IV 22.9 Underweight
```

To **drop** specific columns, use `-` before the column name:

```r
colnames(clinical_data)
```

```
 [1] "patient_id"        "site"              "visit_date"
 [4] "visit_year_month"  "visit_year"        "visit_month"
 [7] "visit_day"         "age"               "sex"
[10] "smoking"           "treatment"         "stage"
[13] "weight_kg"         "height_cm"         "systolic_bp_mmhg"
[16] "diastolic_bp_mmhg" "glucose_mgdl"      "cholesterol_mgdl"
[19] "crp_ngml"          "height_m"          "BMI"
[22] "BMI_class"
```

```r
# drop the helper/intermediate columns we created
clinical_data <- clinical_data |>
  select(-visit_year_month)

colnames(clinical_data)
```

```
 [1] "patient_id"        "site"              "visit_date"
 [4] "visit_year"        "visit_month"       "visit_day"
 [7] "age"               "sex"               "smoking"
[10] "treatment"         "stage"             "weight_kg"
[13] "height_cm"         "systolic_bp_mmhg"  "diastolic_bp_mmhg"
[16] "glucose_mgdl"      "cholesterol_mgdl"  "crp_ngml"
[19] "height_m"          "BMI"               "BMI_class"
```

`select` also supports handy helper functions:

```r
# select all columns that start with "diag"
clinical_data |> select(starts_with("visit")) |> head()
```

```
       visit_date visit_year visit_month visit_day
PT-001 2023-07-20       2023          07        20
PT-002 2020-08-15       2020          08        15
PT-003 2023-10-15       2023          10        15
PT-004 2020-03-06       2020          03        06
PT-005 2023-03-30       2023          03        30
PT-006 2023-04-15       2023          04        15
```

```r
# select all columns that contain "bp"
clinical_data |> select(contains("bp")) |> head()
```

```
       systolic_bp_mmhg diastolic_bp_mmhg
PT-001              164               101
PT-002              143               102
PT-003              152                98
PT-004              142                88
PT-005              160                96
PT-006              150               107
```

2.3.3 Boolean Operators

A **boolean** (or logical) is a data type that can only take two values: `TRUE` or `FALSE`. It is the backbone of filtering and conditional logic in R.

A boolean is the answer to a yes/no question.

For example: *"Is this patient in stage III or IV?"* → `TRUE` or `FALSE`.

2.3.3.1 Comparison operators

These operators compare two values and return a boolean:

Operator	Meaning	Example	Result
`==`	equal to	`3 == 3`	`TRUE`
`!=`	not equal to	`3 != 4`	`TRUE`
`>`	greater than	`5 > 3`	`TRUE`
`<`	less than	`2 < 1`	`FALSE`
`>=`	greater or equal	`3 >= 3`	`TRUE`
`<=`	less or equal	`2 <= 1`	`FALSE`

```r
# simple examples
print(10 > 5)
```

```
[1] TRUE
```

```r
print(10 == 10)
```

```
[1] TRUE
```

```r
print(10 != 10)
```

```
[1] FALSE
```

```r
print(3.14 >= 3.14)
```

```
[1] TRUE
```

When applied to a **vector**, the comparison is done **element-wise** and returns a boolean vector:

```r
ages <- clinical_data$age
head(ages)
```

```
[1] 40 40 59 46 58 53
```

```r
is_over_60 <- ages > 60
head(is_over_60)
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE
```

Under the hood, TRUE is stored as a 1, while FALSE as 0. This is useful when we want to get the number of TRUE values of a vector. How? By summing up all the values:

```r
sum(is_over_60)
```

```
[1] 84
```

2.3.3.2 Logical operators

We often need to combine multiple conditions. This is done with **logical operators**:

Operator	Meaning	Example
`&`	AND — both conditions must be true	`age > 50 & sex == "Female"`
`|`	OR — at least one condition must be true	`stage == "III" | stage == "IV"`
`!`	NOT — negates a condition	`!is.na(crp_ngml)`

```r
# patients who are Female AND over 60
female_over60 <- (clinical_data$sex == "Female") & (clinical_data$age > 60)
sum(female_over60)
```

```
[1] 49
```

```r
# patients in advanced stage (III or IV)
advanced_stage <- (clinical_data$stage == "III") | (clinical_data$stage == "IV")
sum(advanced_stage)
```

```
[1] 91
```

```r
# NOT: patients who are NOT in the Control group
not_control <- clinical_data$treatment != "Control"
sum(not_control)
```

```
[1] 166
```

```r
# equivalent using !
not_control_2 <- !(clinical_data$treatment == "Control")
sum(not_control_2)
```

```
[1] 166
```

A very useful operator when checking membership in a list of values is `%in%`:

```r
# patients from Milan or Rome (equivalent to == "Milan" | == "Rome")
big_sites <- clinical_data$site %in% c("Milan", "Rome")
sum(big_sites)
```

```
[1] 173
```

::: .{exercise #bool-sum}

How many patients have been diagnosed as stage I, in Turin in 2018?

:::

Solution

```r
stageI_turin_2018 <- (clinical_data$site == "Turin") & (clinical_data$stage == "I") & (clinical_d
sum(stageI_turin_2018)
```

```
[1] 4
```

2.3.4 filter and filter out

`filter` keeps only the **rows** that satisfy one or more conditions:

```r
# Keep only patients treated with Drug_A
drug_a_patients <- clinical_data |>
  filter(treatment == "Drug_A")

nrow(drug_a_patients)
```

```
[1] 88
```

```r
head(drug_a_patients)
```

```
       patient_id  site visit_date visit_year visit_month visit_day age    sex
PT-001     PT-001 Milan 2023-07-20       2023          07        20  40   Male
PT-003     PT-003  Rome 2023-10-15       2023          10        15  59 Female
PT-004     PT-004 Turin 2020-03-06       2020          03        06  46   Male
PT-007     PT-007  Rome 2020-09-09       2020          09        09  51   Male
PT-010     PT-010  Rome 2023-09-27       2023          09        27  60   Male
PT-012     PT-012  Rome 2025-03-20       2025          03        20  52   Male
       smoking treatment stage weight_kg height_cm systolic_bp_mmhg
PT-001   never    Drug_A    II      85.9     172.7              164
PT-003  former    Drug_A     I      65.0     165.6              152
PT-004 current    Drug_A     I      95.9     181.5              142
PT-007  former    Drug_A     I      69.7     169.7              137
PT-010  former    Drug_A    II      93.4     186.3              159
PT-012   never    Drug_A   III      84.3     168.8              147
       diastolic_bp_mmhg glucose_mgdl cholesterol_mgdl crp_ngml height_m  BMI
PT-001               101          113              172     5.89    1.727 28.8
PT-003                98          110              146     7.81    1.656 23.7
PT-004                88          120              232     3.35    1.815 29.1
PT-007                96          108              163     5.41    1.697 24.2
PT-010               104          110              185     5.30    1.863 26.9
PT-012                96           86              222    10.45    1.688 29.6
         BMI_class
PT-001  Overweight
PT-003 Underweight
PT-004  Overweight
PT-007 Underweight
PT-010  Overweight
PT-012  Overweight
```

```r
# Patients in advanced stage AND not in control group
advanced_treated <- clinical_data |>
  filter((stage == "III" | stage == "IV") & treatment != "Control")

nrow(advanced_treated)
```

```
[1] 76
```

```r
# Using %in% for multiple values
milan_or_rome <- clinical_data |>
  filter(site %in% c("Milan", "Rome"))

nrow(milan_or_rome)
```

```
[1] 173
```

To **filter out** (i.e., *exclude* rows matching a condition), use `filter_out` (for new versions of R and tidyverse):

```r
# Exclude patients from Naples
no_milan <- clinical_data |>
  filter_out(site %in% c("Milan"))

nrow(no_milan)
```

```
[1] 144
```

If it doesn't work, you could use:

```r
no_milan = clinical_data %>%
    filter(!(site == "Milan"))

nrow(no_milan)
```

```
[1] 144
```

Pipes shine when you chain multiple operations:

```r
# Female patients, over 55, from Milan, keep only key columns
subset_females <- clinical_data |>
  filter(sex == "Female", age > 55, site == "Milan") |>
  select(patient_id, age, stage, treatment, BMI, glucose_mgdl)

head(subset_females)
```

```
       patient_id age stage treatment  BMI glucose_mgdl
PT-011     PT-011  65    II   Control 30.5           94
PT-014     PT-014  71     I   Control 21.5           86
PT-018     PT-018  72   III    Drug_A 25.2           83
PT-037     PT-037  58    II    Drug_B 26.8          108
PT-043     PT-043  60   III    Drug_B 25.9          109
PT-045     PT-045  65    II   Control 27.8           97
```

```r
nrow(subset_females)
```

```
[1] 27
```

Note: inside `filter`, multiple conditions separated by commas are treated as AND.

Exercise 2.2. From the full dataset, extract patients who:

- are in stage II or III,
- have a glucose level between 90 and 130 mg/dL (inclusive),
- are assigned to Drug_A or Drug_B.

Keep only `patient_id`, `site`, `stage`, `treatment`, and `glucose_mgdl`.

Solution

```r
exercise_result <- clinical_data |>
  filter(stage %in% c("II", "III"),
         glucose_mgdl >= 90,
         glucose_mgdl <= 130,
         treatment %in% c("Drug_A", "Drug_B")) |>
  select(patient_id, site, stage, treatment, glucose_mgdl)

head(exercise_result)
```

```
       patient_id  site stage treatment glucose_mgdl
PT-001     PT-001 Milan    II    Drug_A          113
PT-009     PT-009 Milan   III    Drug_B          106
PT-010     PT-010  Rome    II    Drug_A          110
PT-016     PT-016 Turin    II    Drug_B           96
PT-020     PT-020 Turin   III    Drug_B           98
PT-026     PT-026 Turin   III    Drug_A          122
```

```r
nrow(exercise_result)
```

```
[1] 68
```

2.3.5 group_by, summarize and reframe

One of the most powerful operations in data analysis is **grouping** data and computing summaries per group. In R, this is done with `group_by` + `summarize`.

How it works:

1. `group_by(column)` tells R: *"split the dataframe into groups based on the values of this column"*
2. `summarize(...)` computes one summary value per group

```
Original data:                 After group_by(treatment):

patient  treatment  glucose       GROUP: Control
PT-001   Control    105           PT-001  Control  105
PT-002   Drug_A     118    --->   PT-007  Control  98
PT-003   Drug_B     97            ...
PT-007   Control    98     --->   GROUP: Drug_A
...                               PT-002  Drug_A  118
                                  ...
```

We can perform several operation with this combo, and answer many questions. For example "*Which is the mean and sd values of glucose in each treatment group? How many patients we have per treatment group?*".

```r
# Mean glucose per treatment group
glucose_by_treatment <- clinical_data |>
  group_by(treatment) |>
  summarize(mean_glucose = mean(glucose_mgdl),
            sd_glucose   = sd(glucose_mgdl),
            n_patients   = n())

glucose_by_treatment
```

```
# A tibble: 3 x 4
  treatment mean_glucose sd_glucose n_patients
  <fct>            <dbl>      <dbl>      <int>
1 Control           109.       15.9         74
2 Drug_A            110.       14.9         88
3 Drug_B            108.       16.0         78
```

You can also group by multiple columns at once.

"*For each treatment/sex combination, which is the mean BMI and the median of the age?*"

```r
# Summary by treatment AND sex
summary_by_group <- clinical_data |>
  group_by(treatment, sex) |>
  summarize(mean_BMI    = round(mean(BMI), 2),
            median_age  = median(age),
            n           = n(),
```

```
            .groups = "drop")   # good practice: drop grouping after summarize

summary_by_group
```

```
# A tibble: 6 x 5
  treatment sex    mean_BMI median_age     n
  <fct>     <fct>     <dbl>      <dbl> <int>
1 Control   Female     27.1       59      39
2 Control   Male       28.5       56      35
3 Drug_A    Female     25.6       58      43
4 Drug_A    Male       27.2       55      45
5 Drug_B    Female     25.2       59.5    36
6 Drug_B    Male       26.8       56      42
```

`summarize` returns **one row per group**. If you need to return **multiple rows per group** (e.g., all quantiles), use `reframe` instead:

```
# Get quartiles of CRP per treatment group
crp_quantiles <- clinical_data |>
  group_by(treatment) |>
  reframe(quantile_value = quantile(crp_ngml, probs = c(0.25, 0.5, 0.75)),
          quantile_name  = c("Q1", "median", "Q3"))

crp_quantiles
```

```
# A tibble: 9 x 3
  treatment quantile_value quantile_name
  <fct>              <dbl> <chr>
1 Control             5.18 Q1
2 Control             8.39 median
3 Control            13.8  Q3
4 Drug_A              5.25 Q1
5 Drug_A              7.3  median
6 Drug_A             11.4  Q3
7 Drug_B              5.36 Q1
8 Drug_B              8.52 median
9 Drug_B             12.7  Q3
```

Another useful function: `count`, which is a shortcut for `group_by` + `summarize(n = n())`:

```
# Count patients per site and stage
clinical_data |>
  count(site, stage, sort = TRUE)
```

```
    site stage  n
1  Milan    II 35
2   Rome     I 33
3  Milan     I 23
4  Turin    II 22
5   Rome   III 21
6  Turin     I 20
7  Milan   III 19
8  Milan    IV 19
9   Rome    II 16
10 Turin   III 13
11 Turin    IV 12
12  Rome    IV  7
```

Exercise 2.3. For each combination of `site` and `sex`, compute: - the number of patients, - the mean systolic blood pressure, - the proportion of hypertensive patients (hint: hypertension is when systolic blood pression is > 140).

Solution

```r
bp_summary <- clinical_data |>
    mutate(hypertension = systolic_bp_mmhg > 140) |>
    group_by(site, sex) |>
    summarize(
        n_patients          = n(),
        mean_systolic_bp_mmhg    = round(mean(systolic_bp_mmhg), 1),
        prop_hypertensive   = round(mean(hypertension), 3),
        .groups = "drop"
    )

bp_summary
```

```
# A tibble: 6 x 5
  site  sex    n_patients mean_systolic_bp_mmhg prop_hypertensive
  <fct> <fct>       <int>                 <dbl>             <dbl>
1 Milan Female         48                  146.             0.708
2 Milan Male           48                  149.             0.75
3 Rome  Female         39                  144.             0.667
4 Rome  Male           38                  148.             0.763
5 Turin Female         31                  143.             0.581
6 Turin Male           36                  150.             0.889
```

2.3.6 Join / Merge

In real-world analyses, data about the same subjects often lives in **separate files**: demographic data here, lab results there, outcomes somewhere else. To combine them, we **join** them on a shared key column (usually a patient ID).

Let's load the separate lab results file:

```
lab_data <- read.csv("data/lab_results.csv")
str(lab_data)
```

```
'data.frame':   160 obs. of  6 variables:
 $ patient_id     : chr  "PT-002" "PT-003" "PT-004" "PT-005" ...
 $ hemoglobin_gdl : num  15 11.8 16 15.4 14 16.2 10.5 15.6 12.4 14.5 ...
 $ wbc_10e9l      : num  7.49 7.46 7.09 6.63 6.99 9.26 4.79 5.65 8.39 7.87 ...
 $ platelets_10e9l: int  270 252 233 192 195 188 163 269 165 263 ...
 $ alt_ul         : int  56 27 20 16 61 51 50 18 47 51 ...
 $ creatinine_mgdl: num  1.58 1.09 1.17 1.06 1.15 1.06 1.21 1.08 1.12 1.03 ...
```

```
head(lab_data)
```

```
  patient_id hemoglobin_gdl wbc_10e9l platelets_10e9l alt_ul creatinine_mgdl
1     PT-002           15.0      7.49             270     56            1.58
2     PT-003           11.8      7.46             252     27            1.09
3     PT-004           16.0      7.09             233     20            1.17
4     PT-005           15.4      6.63             192     16            1.06
5     PT-007           14.0      6.99             195     61            1.15
6     PT-008           16.2      9.26             188     51            1.06
```

The key column is `patient_id`. Not all patients have lab results. Depending on how we want to handle this, we choose different types of join:

Join type	What it does
`left_join(x, y)`	Keep **all rows from x**, add y columns where available (NA if no match)
`right_join(x, y)`	Keep all rows from y
`inner_join(x, y)`	Keep only rows that exist in **both** x and y
`full_join(x, y)`	Keep all rows from both, NA where no match

Let's see all approaches:

```r
# left_join: keep all patients, add lab data where available
left_df <- clinical_data |>
  left_join(lab_data, by = "patient_id")

nrow(left_df)
```

```
[1] 240
```

```r
# patients without lab data will have NA in lab columns
sum(is.na(left_df$hemoglobin_gdl))
```

```
[1] 80
```

Function `is.na()` returns TRUE when a value is NA, and FALSE if a value is an acutal value. Here an example:

```r
v1 = c(1, NA, 4) # create a vector with one NA

is.na(v1) # returns a vector with TRUE in position 2, and FALSE elsewhere
```

```
[1] FALSE  TRUE FALSE
```

```r
summary(left_df)
```

```
  patient_id           site     visit_date         visit_year
 Length:240         Milan:96   Length:240         Length:240
 Class :character   Rome :77   Class :character   Class :character
 Mode  :character   Turin:67   Mode  :character   Mode  :character




 visit_month         visit_day              age           sex         smoking
 Length:240         Length:240         Min.   :26.0   Female:118   current:63
 Class :character   Class :character   1st Qu.:49.0   Male  :122   former :90
 Mode  :character   Mode  :character   Median :57.5                never  :87
                                       Mean   :57.3
                                       3rd Qu.:65.0
                                       Max.   :85.0

   treatment  stage      weight_kg        height_cm     systolic_bp_mmhg
 Control:74   I  :76   Min.   : 39.40   Min.   :148.9   Min.   :115.0
 Drug_A :88   II :73   1st Qu.: 65.35   1st Qu.:162.2   1st Qu.:139.0
```

```
 Drug_B :78   III:53   Median : 75.65   Median :168.5   Median :146.5
              IV :38   Mean   : 76.87   Mean   :169.2   Mean   :146.7
                       3rd Qu.: 87.83   3rd Qu.:174.8   3rd Qu.:153.2
                       Max.   :132.90   Max.   :200.1   Max.   :192.0

 diastolic_bp_mmhg  glucose_mgdl   cholesterol_mgdl    crp_ngml
 Min.   : 73.00    Min.   : 60.0   Min.   :113.0    Min.   : 1.92
 1st Qu.: 90.75    1st Qu.: 98.0   1st Qu.:189.8    1st Qu.: 5.30
 Median : 97.00    Median :109.0   Median :207.0    Median : 8.37
 Mean   : 96.11    Mean   :108.8   Mean   :207.9    Mean   :10.35
 3rd Qu.:102.00    3rd Qu.:118.0   3rd Qu.:226.0    3rd Qu.:13.02
 Max.   :120.00    Max.   :149.0   Max.   :307.0    Max.   :42.09

    height_m          BMI              BMI_class   hemoglobin_gdl
 Min.   :1.489   Min.   :16.00   Underweight: 67   Min.   : 9.20
 1st Qu.:1.623   1st Qu.:23.77   Normal     : 58   1st Qu.:12.40
 Median :1.685   Median :26.85   Overweight :102   Median :13.25
 Mean   :1.692   Mean   :26.71   Obese      : 13   Mean   :13.50
 3rd Qu.:1.748   3rd Qu.:29.80                     3rd Qu.:14.70
 Max.   :2.001   Max.   :40.70                     Max.   :17.00
                                                   NA's   :80
   wbc_10e9l      platelets_10e9l     alt_ul       creatinine_mgdl
 Min.   : 2.310   Min.   : 60.0   Min.   : 10.00   Min.   :0.630
 1st Qu.: 5.518   1st Qu.:191.8   1st Qu.: 32.75   1st Qu.:1.018
 Median : 7.000   Median :232.0   Median : 46.50   Median :1.140
 Mean   : 6.758   Mean   :231.7   Mean   : 52.20   Mean   :1.138
 3rd Qu.: 7.902   3rd Qu.:271.0   3rd Qu.: 61.00   3rd Qu.:1.252
 Max.   :11.380   Max.   :391.0   Max.   :161.00   Max.   :1.580
 NA's   :80       NA's   :80      NA's   :80       NA's   :80
```

```r
# inner_join: keep only the 90 patients who have both clinical and lab data
matched_only <- clinical_data |>
  inner_join(lab_data, by = "patient_id")

nrow(matched_only)
```

```
[1] 160
```

In most clinical research scenarios, `left_join` is the most common choice: you start with your full patient list and enrich it with additional data, accepting that some patients may lack certain measurements (NAs).

Exercise 2.4.

1. Perform a `left_join` of `clinical_data` with `lab_data`.

2. Among patients who have lab data available, compute the mean hemoglobin and mean creatinine per treatment group.
3. Which group has the highest mean creatinine?

Solution

```
# 1. Left join
joined_df <- clinical_data |>
  left_join(lab_data, by = "patient_id")

# 2. Summarize (filter out NA first)
lab_summary <- joined_df |>
  filter(!is.na(hemoglobin_gdl)) |> # filter for what is NOT NA in hemoglobin
  group_by(treatment) |>
  summarize(mean_hemoglobin   = round(mean(hemoglobin_gdl), 2),
            mean_creatinine   = round(mean(creatinine_mgdl), 3),
            n                 = n(),
            .groups = "drop")

lab_summary
```

```
# A tibble: 3 x 4
  treatment mean_hemoglobin mean_creatinine     n
  <fct>               <dbl>           <dbl> <int>
1 Control              13.7            1.16    51
2 Drug_A               13.5            1.13    62
3 Drug_B               13.2            1.13    47
```

```
# 3. Highest mean creatinine
lab_summary |> arrange(desc(mean_creatinine))
```

```
# A tibble: 3 x 4
  treatment mean_hemoglobin mean_creatinine     n
  <fct>               <dbl>           <dbl> <int>
1 Control              13.7            1.16    51
2 Drug_B               13.2            1.13    47
3 Drug_A               13.5            1.13    62
```

2.4 Save your results

After all this work, let's save the enriched and cleaned dataset:

```r
# Save the left-joined dataset with all derived variables
final_data <- clinical_data |>
  left_join(lab_data, by = "patient_id")

write.csv(final_data,
          file = "output/clinical_data_enriched.csv",
          row.names = FALSE)
```

2.5 Home exercise

Try to do the following on your own:

1. Load `clinical_data_enriched.csv` from the output folder.
2. Clean up: keep only the columns `patient_id`, `site`, `sex`, `age`, `treatment`, `stage`, `BMI`, `BMI_class`, `glucose_mgdl`, `cholesterol_mgdl`, `crp_ngml`, `hemoglobin_gdl`, `creatinine_mgdl`.
3. Restrict to patients who have complete data (no NAs) — hint: use function `drop_na()`, see this.
4. For each stage, compute: number of patients, mean age, mean BMI, and proportion of obese patients.
5. Save the summary table to `output/summary_by_stage.csv`.

```r
library(tidyverse)

# 1. Load the enriched dataset
clinical_data_enriched <- read.csv("output/clinical_data_enriched.csv")

# 2. Keep only the columns of interest
clean_data <- clinical_data_enriched |>
  select(patient_id, site, sex, age, treatment, stage,
         BMI, BMI_class, glucose_mgdl, cholesterol_mgdl,
         crp_ngml, hemoglobin_gdl, creatinine_mgdl)

# 3. Restrict to patients with complete data (no NAs)
clean_data <- clean_data |>
  drop_na()

cat("Patients with complete data:", nrow(clean_data), "\n")
```

```
Patients with complete data: 160
```

```r
# 4. Summary by stage
summary_by_stage <- clean_data |>
  group_by(stage) |>
  summarize(
    n_patients       = n(),
    mean_age         = round(mean(age), 1),
    mean_BMI         = round(mean(BMI), 2),
    prop_obese       = round(mean(BMI_class == "Obese"), 3),
    .groups = "drop"
  )

print(summary_by_stage)
```

```
# A tibble: 4 x 5
  stage n_patients mean_age mean_BMI prop_obese
  <chr>      <int>    <dbl>    <dbl>      <dbl>
1 I             47     55.2     27.2      0.043
2 II            53     57.1     27.3      0.038
3 III           35     56.2     28.2      0.029
4 IV            25     57.5     24.2      0.16
```

```r
# 5. Save
write.csv(summary_by_stage,
          file = "output/summary_by_stage.csv",
          row.names = FALSE)
```

2.6 Bonus: how the clinical data were created

If you want to understand how we generated the dataset for this lesson, download the R script Notice the use of `set.seed` to make results reproducible.

Chapter 3

Data visualization

Data visualization is a crucial step of data analysis, as it can help us in:

- Understanding the distribution of our data
- Evaluate outliers
- Visualize immediately the differences between groups
- Communicate our results to peers

I suggest you to visualize data as soon as you have finished to prepare them, prior to any statistical test (even prior to test for parametric assumptions), that's why today we are doing it.

3.1 Data visualization in R with ggplot2

`ggplot2` is a package built for data visualization in R. It is simple to use and it works mainly on dataframes.

To create a plot, we need to use different functions, that act as single building blocks to change single aspects of the plot. Here an example of mock code to create a plot with all the functions:

```
ggplot() +
  geom_<type>() +
  labs() +
  scale_<type>() +
  facet_<type>() +
  coord_<type>() +
  theme_<type>
```

But, do we need all of them? No, the basic plots can be created with just `ggplot() + geom_<type>()`.

3.2 Geom

The `geom` function specifies the "geometry" we want to give to the data: basically, which type of graph we want (barplot, scatter, histogram, etc). It is the necessary and sufficient function to create a plot.

Here the list of all possible geom to use with the basic package.

First of all, let's load the data from previous session, and keep only the rows with all the information:

```
suppressPackageStartupMessages(library(tidyverse))
df <- read.csv("output/clinical_data_enriched.csv", header = T, stringsAsFactors = T)
df <- na.omit(df)
str(df)
```

```
'data.frame':   160 obs. of  26 variables:
 $ patient_id       : Factor w/ 240 levels "PT-001","PT-002",..: 2 3 4 5 7 8 9 10 11 12
 $ site             : Factor w/ 3 levels "Milan","Rome",..: 2 2 3 3 2 2 1 2 1 2 ...
 $ visit_date       : Factor w/ 218 levels "2018-01-14","2018-01-24",..: 75 150 52 104
 $ visit_year       : int  2020 2023 2020 2023 2020 2025 2020 2023 2023 2025 ...
 $ visit_month      : int  8 10 3 3 9 10 4 9 6 3 ...
 $ visit_day        : int  15 15 6 30 9 7 19 27 29 20 ...
 $ age              : int  40 59 46 58 51 34 43 60 65 52 ...
 $ sex              : Factor w/ 2 levels "Female","Male": 2 1 2 2 2 1 2 2 1 2 ...
 $ smoking          : Factor w/ 3 levels "current","former",..: 1 2 1 2 2 3 3 2 1 3 ..
 $ treatment        : Factor w/ 3 levels "Control","Drug_A",..: 3 2 2 3 2 1 3 2 1 2 ..
 $ stage            : Factor w/ 4 levels "I","II","III",..: 1 1 1 3 1 1 3 2 2 3 ...
 $ weight_kg        : num  100.9 65 95.9 75.7 69.7 ...
 $ height_cm        : num  185 166 182 168 170 ...
 $ systolic_bp_mmhg : int  143 152 142 160 137 146 151 159 125 147 ...
 $ diastolic_bp_mmhg: int  102 98 88 96 96 106 94 104 76 96 ...
 $ glucose_mgdl     : int  105 110 120 131 108 93 106 110 94 86 ...
 $ cholesterol_mgdl : int  239 146 232 192 163 199 200 185 209 222 ...
 $ crp_ngml         : num  6.13 7.81 3.35 21.72 5.41 ...
 $ height_m         : num  1.85 1.66 1.81 1.68 1.7 ...
 $ BMI              : num  29.4 23.7 29.1 26.8 24.2 29.8 24.8 26.9 30.5 29.6 ...
 $ BMI_class        : Factor w/ 4 levels "Normal","Obese",..: 3 4 3 3 4 3 4 3 1 3 ...
 $ hemoglobin_gdl   : num  15 11.8 16 15.4 14 16.2 10.5 15.6 12.4 14.5 ...
 $ wbc_10e9l        : num  7.49 7.46 7.09 6.63 6.99 9.26 4.79 5.65 8.39 7.87 ...
 $ platelets_10e9l  : int  270 252 233 192 195 188 163 269 165 263 ...
 $ alt_ul           : int  56 27 20 16 61 51 50 18 47 51 ...
```

```
 $ creatinine_mgdl  : num  1.58 1.09 1.17 1.06 1.15 1.06 1.21 1.08 1.12 1.03 ...
 - attr(*, "na.action")= 'omit' Named int [1:80] 1 6 14 21 25 33 46 49 52 54 ...
  ..- attr(*, "names")= chr [1:80] "1" "6" "14" "21" ...
```

```r
summary(df)
```

```
   patient_id     site         visit_date    visit_year    visit_month
 PT-002 :  1   Milan:65   2020-07-21:  3   Min.   :2018   Min.   : 1.000
 PT-003 :  1   Rome :50   2018-03-03:  2   1st Qu.:2020   1st Qu.: 3.750
 PT-004 :  1   Turin:45   2018-07-18:  2   Median :2023   Median : 6.000
 PT-005 :  1              2020-05-23:  2   Mean   :2022   Mean   : 5.969
 PT-007 :  1              2020-10-16:  2   3rd Qu.:2025   3rd Qu.: 8.000
 PT-008 :  1              2023-03-06:  2   Max.   :2025   Max.   :11.000
 (Other):154              (Other)   :147
   visit_day          age            sex        smoking     treatment  stage
 Min.   : 1.00   Min.   :33.00   Female:83   current:44   Control:51   I  :47
 1st Qu.:10.00   1st Qu.:48.00   Male  :77   former :55   Drug_A :62   II :53
 Median :18.00   Median :56.50               never  :61   Drug_B :47   III:35
 Mean   :17.16   Mean   :56.41                                         IV :25
 3rd Qu.:24.00   3rd Qu.:64.00
 Max.   :31.00   Max.   :85.00

   weight_kg        height_cm     systolic_bp_mmhg diastolic_bp_mmhg
 Min.   : 39.40   Min.   :148.9   Min.   :120.0    Min.   : 73.00
 1st Qu.: 66.17   1st Qu.:162.4   1st Qu.:138.0    1st Qu.: 89.00
 Median : 75.75   Median :168.2   Median :146.0    Median : 96.00
 Mean   : 77.59   Mean   :169.2   Mean   :145.9    Mean   : 95.46
 3rd Qu.: 88.33   3rd Qu.:174.6   3rd Qu.:153.0    3rd Qu.:101.00
 Max.   :132.90   Max.   :191.3   Max.   :178.0    Max.   :120.00

  glucose_mgdl cholesterol_mgdl    crp_ngml         height_m
 Min.   : 75   Min.   :113.0    Min.   : 1.930   Min.   :1.489
 1st Qu.: 98   1st Qu.:187.8    1st Qu.: 5.300   1st Qu.:1.624
 Median :109   Median :205.5    Median : 8.465   Median :1.682
 Mean   :110   Mean   :206.1    Mean   :10.525   Mean   :1.692
 3rd Qu.:119   3rd Qu.:223.2    3rd Qu.:13.165   3rd Qu.:1.746
 Max.   :147   Max.   :307.0    Max.   :42.090   Max.   :1.913

      BMI              BMI_class  hemoglobin_gdl    wbc_10e9l
 Min.   :16.00   Normal     :46   Min.   : 9.20   Min.   : 2.310
 1st Qu.:24.18   Obese      : 9   1st Qu.:12.40   1st Qu.: 5.518
 Median :26.90   Overweight :65   Median :13.25   Median : 7.000
 Mean   :26.98   Underweight:40   Mean   :13.50   Mean   : 6.758
 3rd Qu.:30.50                    3rd Qu.:14.70   3rd Qu.: 7.902
 Max.   :40.70                    Max.   :17.00   Max.   :11.380
```

```
platelets_10e9l     alt_ul       creatinine_mgdl
Min.   : 60.0   Min.   : 10.00   Min.   :0.630
1st Qu.:191.8   1st Qu.: 32.75   1st Qu.:1.018
Median :232.0   Median : 46.50   Median :1.140
Mean   :231.7   Mean   : 52.20   Mean   :1.138
3rd Qu.:271.0   3rd Qu.: 61.00   3rd Qu.:1.252
Max.   :391.0   Max.   :161.00   Max.   :1.580
```

3.2.1 Barplot

We create a basic barplot to show the number of patient for each stage using the `geom_bar` function:

```r
suppressPackageStartupMessages(library(ggplot2))
```

```r
ggplot(df) +
  geom_bar(mapping = aes(x = stage))
```

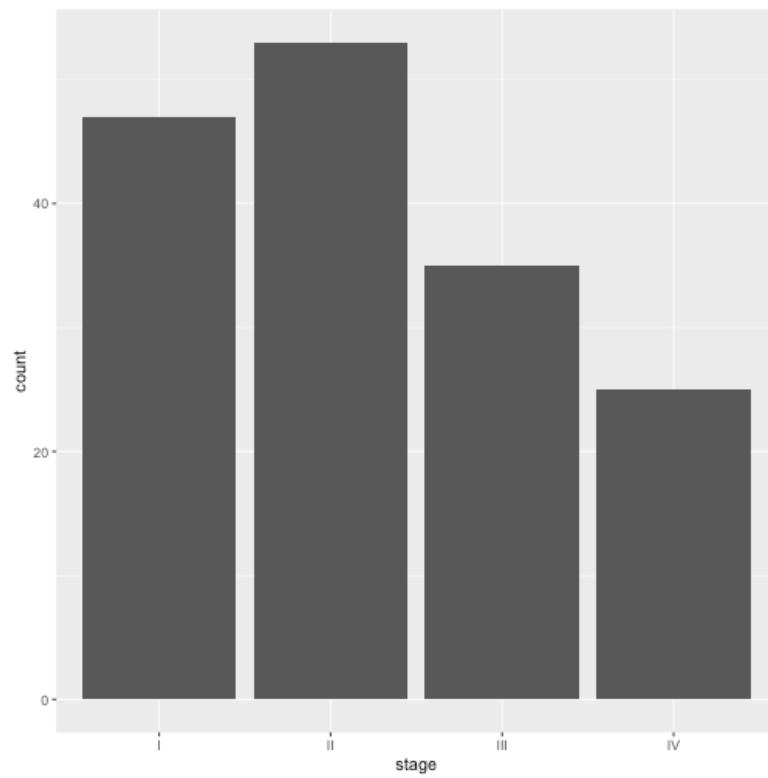

Let's add some colors. We can set colors for all bars:

```
ggplot(df) +
  geom_bar(mapping = aes(x = stage), fill = "blue3")
```

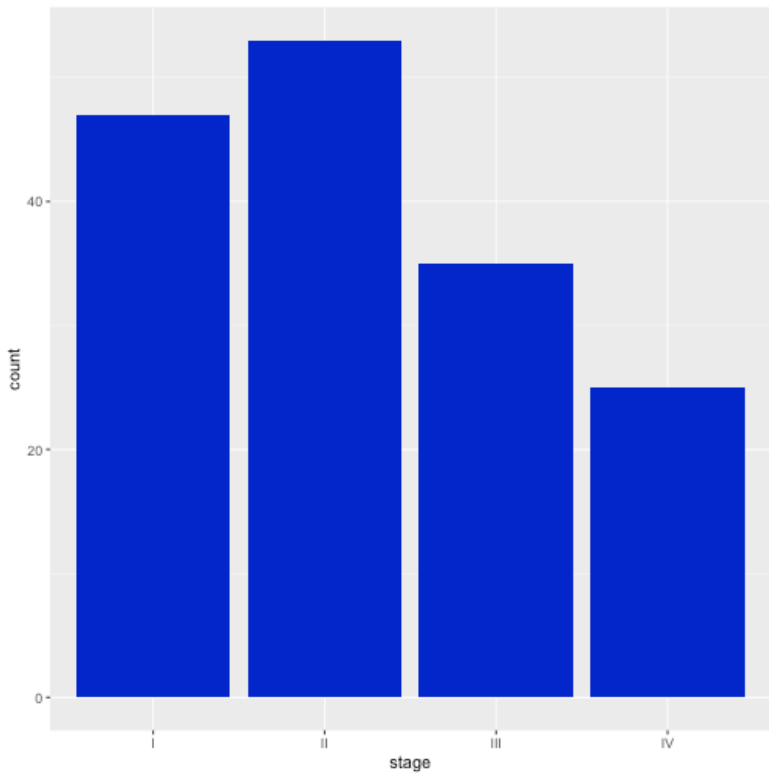

Or based on stage:

```
ggplot(df) +
  geom_bar(mapping = aes(x = stage, fill = stage))
```

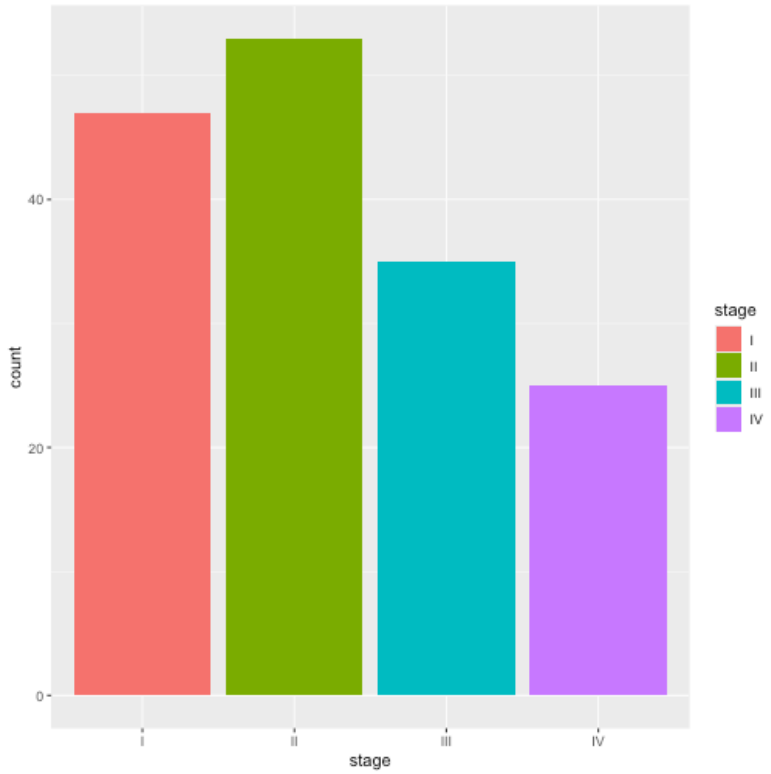

One of the main features of ggplot is that we can easily split data on x axis based on another column. Let's say we want a bar for each treatment within each stage; we can achieve it setting the column in another aes option:

```
ggplot(df) +
  geom_bar(mapping = aes(x = stage, fill = treatment), position = "dodge")
```

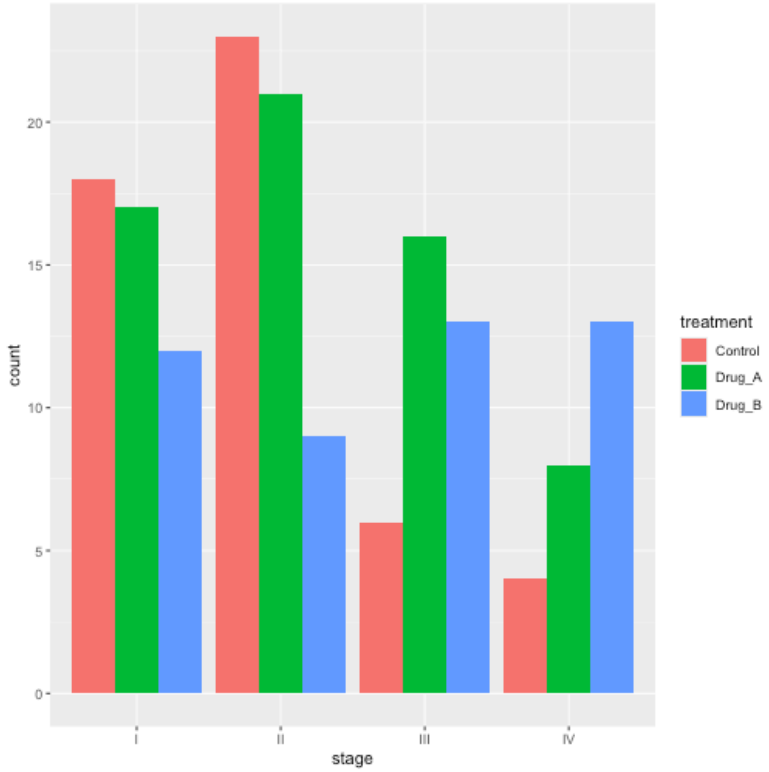

We can set different options in geom bar (both in aes or not), like color (the color of the outline of the bar), linewidth (the size of the outline), linetype (type of the outline of the bars, like dotted or continuous). Here you can find all the options. Here an example:

```
all_n_barplot <- ggplot(df) +
  geom_bar(mapping = aes(x = stage, fill = treatment, linetype = site), position = "dodge", color
all_n_barplot
```

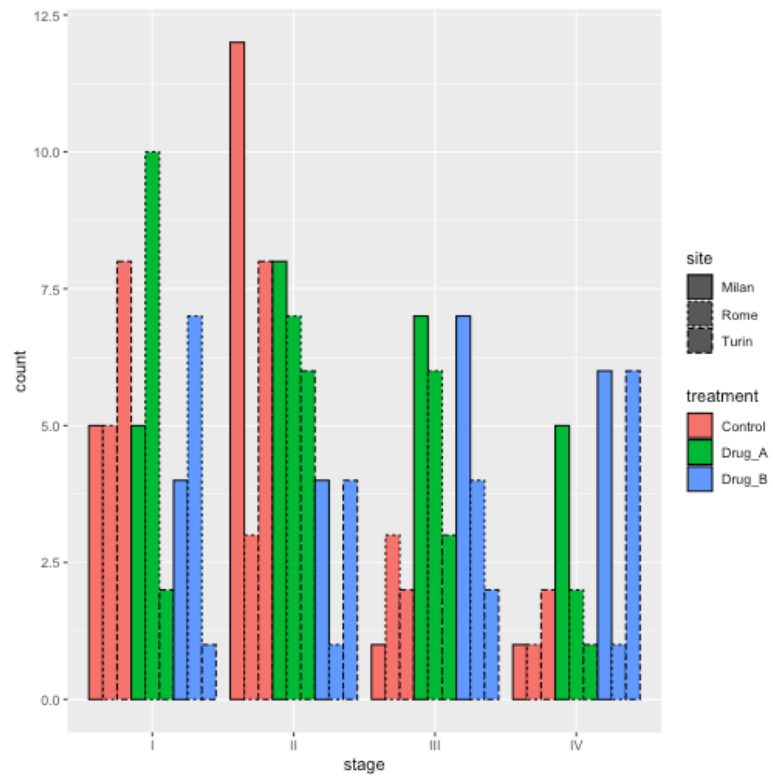

3.2.1.1 Stat identity

In this case, geom_bar counts the occurrence of each label of stage and plot them. But what if we have a dataframe with the actual data we want to plot? Let's say that we want to plot the mean values of "diastolic_bp_mmhg" of Stage I females for each treatment; first, we have to create the dataframe for it by filtering the data and calculating the mean:

```
stageI_femalse_diastolic_means <- df %>%
  filter(stage == "I" & sex == "Female") %>%
  group_by(treatment) %>%
  summarise("mean_diastolic" = mean(diastolic_bp_mmhg))
stageI_femalse_diastolic_means
```

```
# A tibble: 3 x 2
  treatment mean_diastolic
  <fct>              <dbl>
1 Control             96.7
```

```
2 Drug_A              92
3 Drug_B              92
```

Then we can use the function specifying both x and y as aes:

```
stageI_female_diastolic_means_barplot <- ggplot(stageI_femalse_diastolic_means) +
  geom_bar(mapping = aes(x = treatment, y = mean_diastolic, fill = treatment), stat = "identity")
stageI_female_diastolic_means_barplot
```

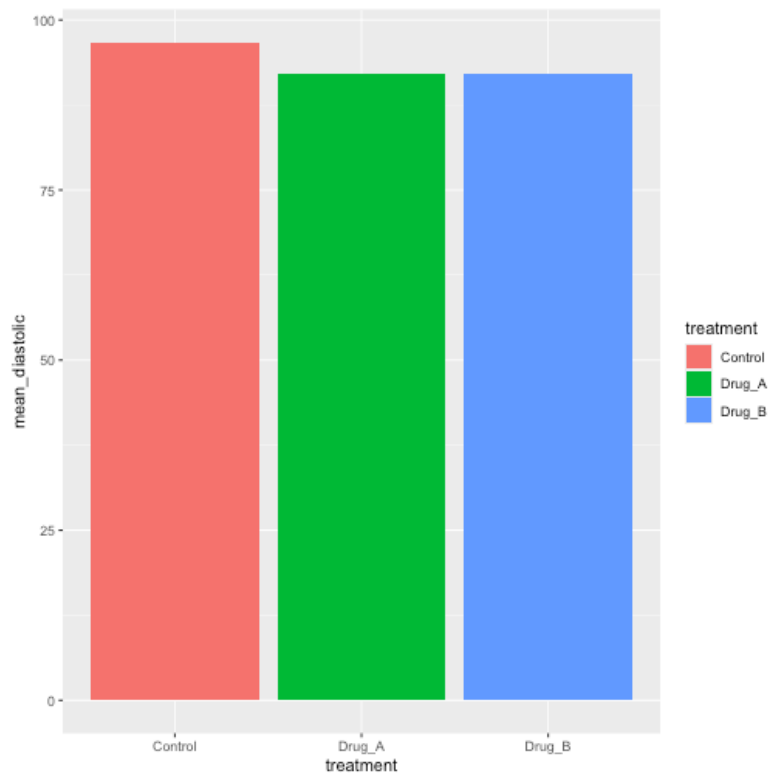

3.2.2 Scatter plot

If we want to create a scatter plot, we can use `geom_point` function.

We can create the plot between weight and height:

```
ggplot(df) +
  geom_point(aes(x = height_cm, y = weight_kg))
```

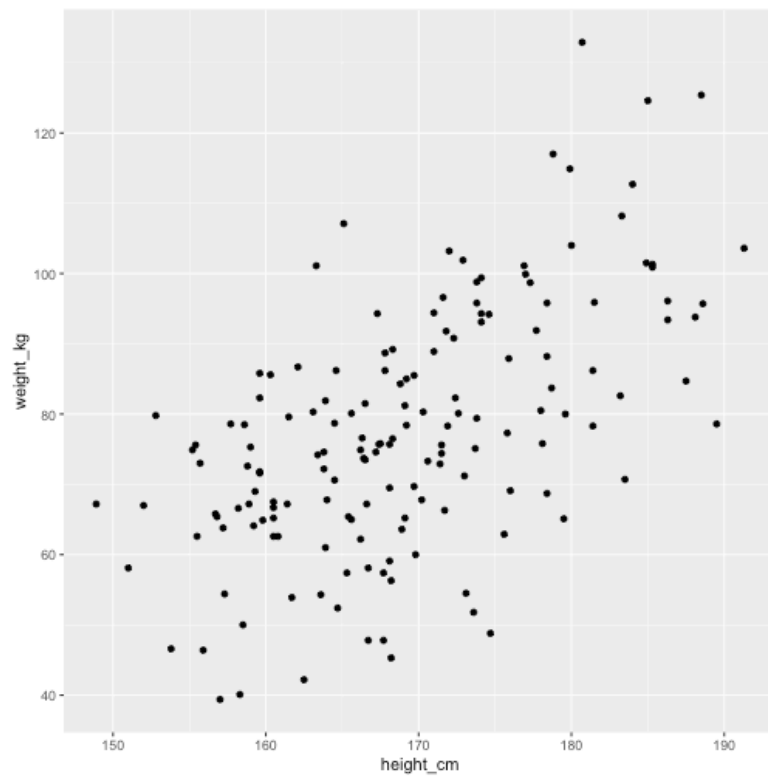

Just from this plot we can see that there is a positive correlation. We can also change color, shape size of the dots. As before, all that is inserted in aes will apply to each dot based on the value of the corresponding column, while outside aes is applied to all dots:

```
corr_plot <- ggplot(df) +
  geom_point(aes(x = height_cm, y = weight_kg, color = age), size = 3, shape = 18)
corr_plot
```

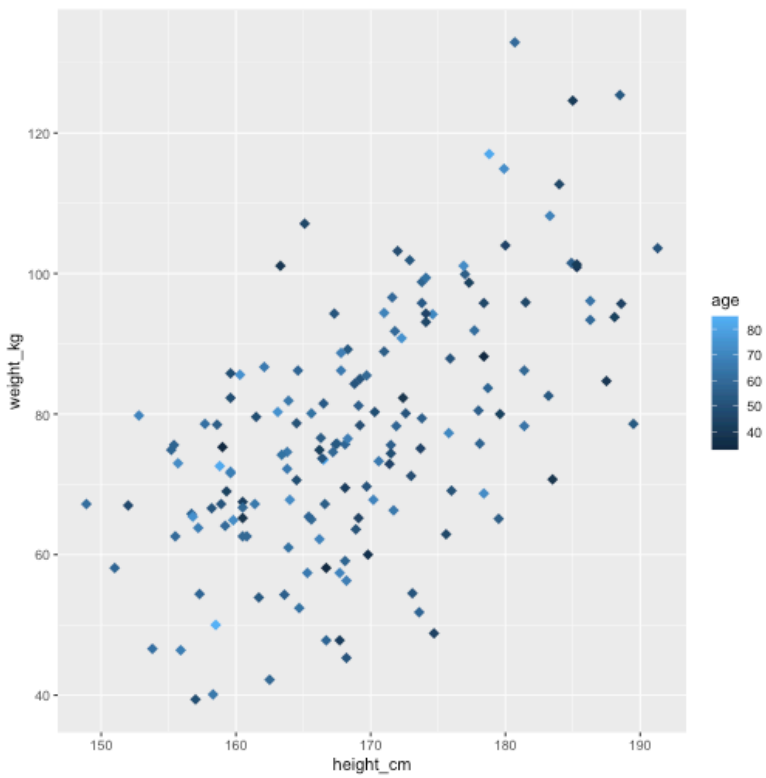

Here the complete list of options for geom_point.

3.2.3 Boxplot

Next, we will see how to create a boxplot, that is one of the most used plot when showing data distribution. Here we look at the diastolic bp in different smoking levels.

```
diastolic_bp_boxplot <- ggplot(df) +
  geom_boxplot(aes(x = smoking, y = diastolic_bp_mmhg, fill = smoking), color = "black", outlier.
diastolic_bp_boxplot
```

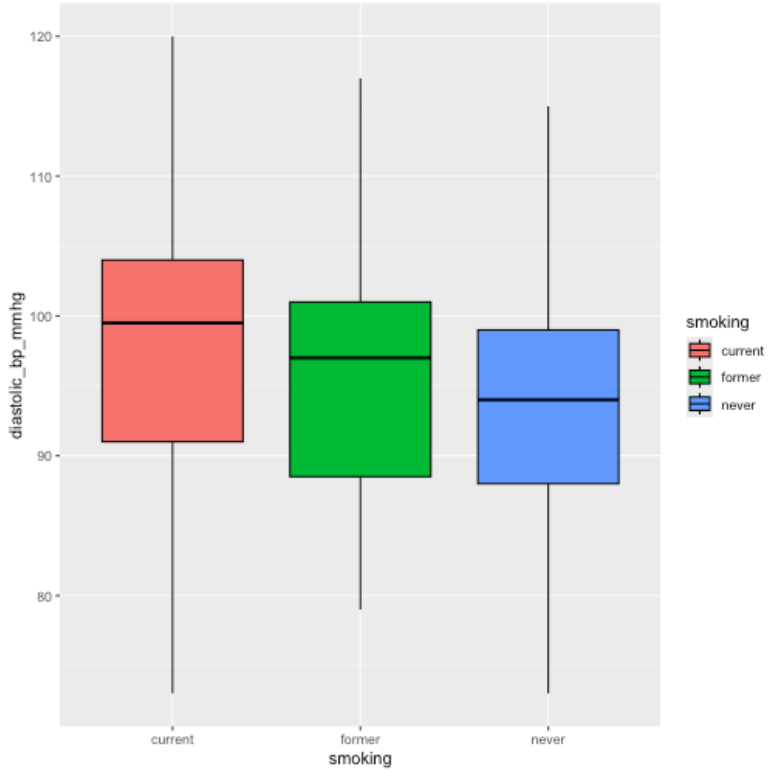

Here the complete list of options for geom_boxplot.

3.2.4 Density plot

Another way to visualize the distribution of the data is the density plot. Let's see the distribution of BMI among stage groups.

```
BMI_stages_density <- ggplot(df) +
  geom_density(aes(x = BMI, fill = stage), alpha = 0.3)
BMI_stages_density
```

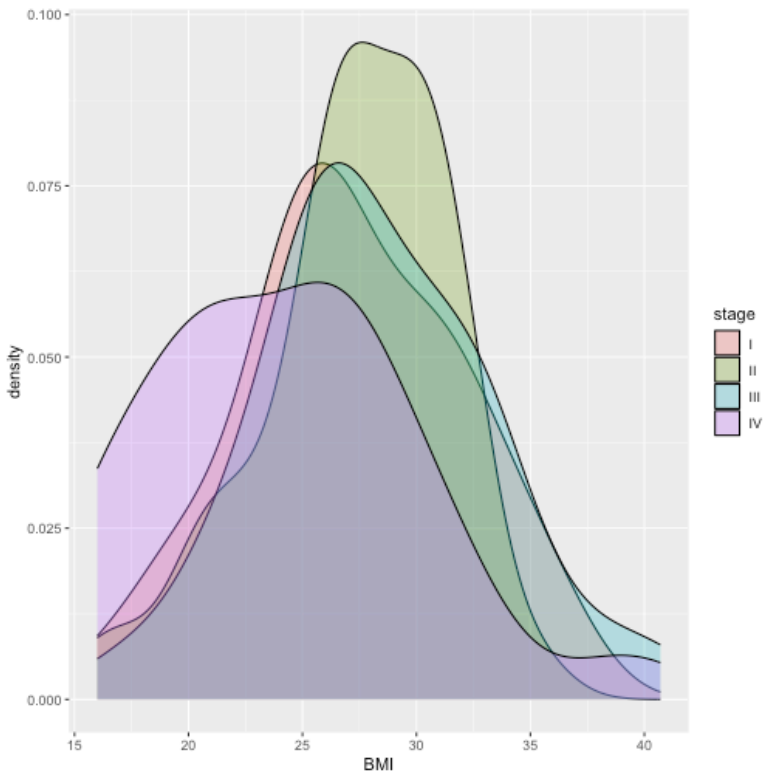

3.2.5 Combine geom

We can also combine geometries. For example, let's add to the *stageI_female_diastolic_means_barplot* the actual values with `geom_text`:

```
stageI_female_diastolic_means_barplot <- stageI_female_diastolic_means_barplot +
  geom_text(aes(x = treatment, y = mean_diastolic, label = round(mean_diastolic)), nudge_y = 5)
stageI_female_diastolic_means_barplot
```

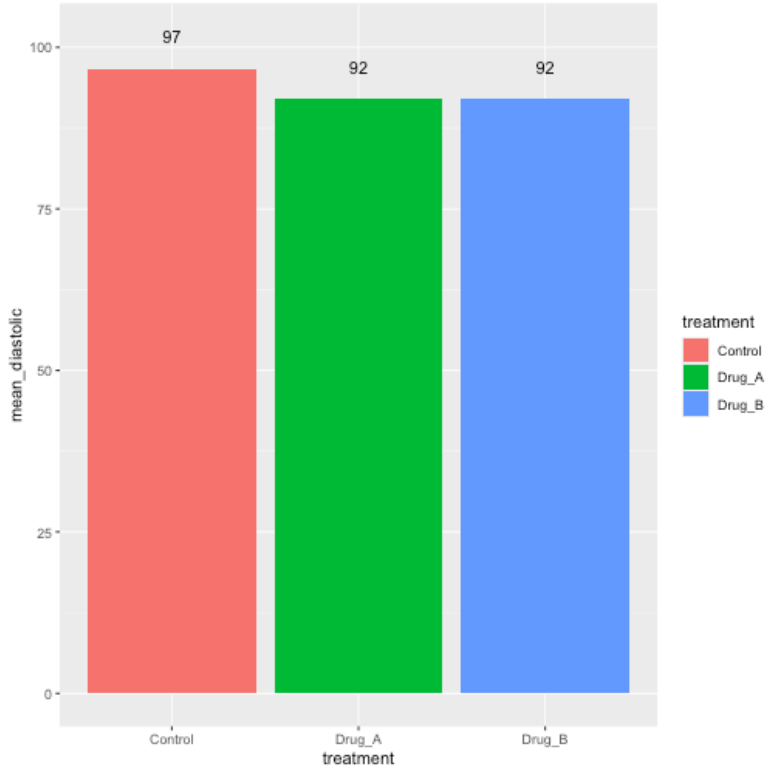

We can combine as many as we want, another example is combine geom_boxplot with geom_point to add individual data points on the plot.

```
ggplot(df) +
  geom_boxplot(aes(x = stage, y = diastolic_bp_mmhg, fill = sex),
               color = "black", outlier.color = "black", position = position_dodge(wid
  geom_jitter(aes(x = stage, y = diastolic_bp_mmhg, fill = sex),
             position = position_dodge(width = 0.75), color = "black", alpha = 0.6, sha
```

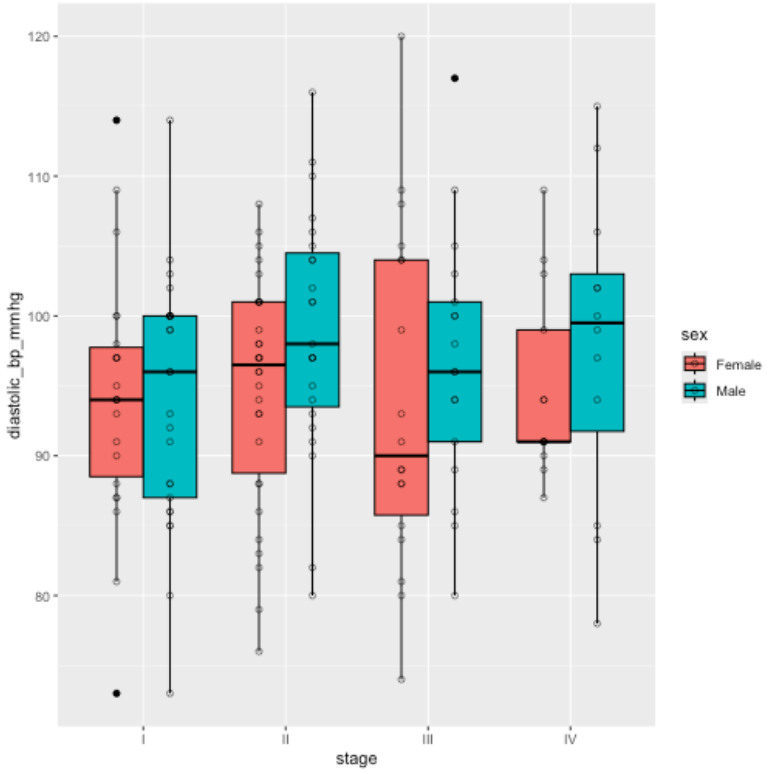

3.3 Labs

All these plots are a good starting point, but we can now start to personalize them. The first thing we can change are labs.

Labs comprehends axis title, plot title and subtitle, legend titles, caption and tag (here the reference page with examples).

```
ggplot() +
  geom_<type>() +
  labs()
```

Let's change the labels of the correlation plot:

```
corr_plot <- corr_plot +
  labs(title = "Correlation between height and weight",
       x = "Height (cm)",
       y = "Weight (kg)",
```

```
        caption = paste("Correlation between height and weight on", dim(df)[1], "patient
                        min(df$age), "-", max(df$age), "years"))
corr_plot
```

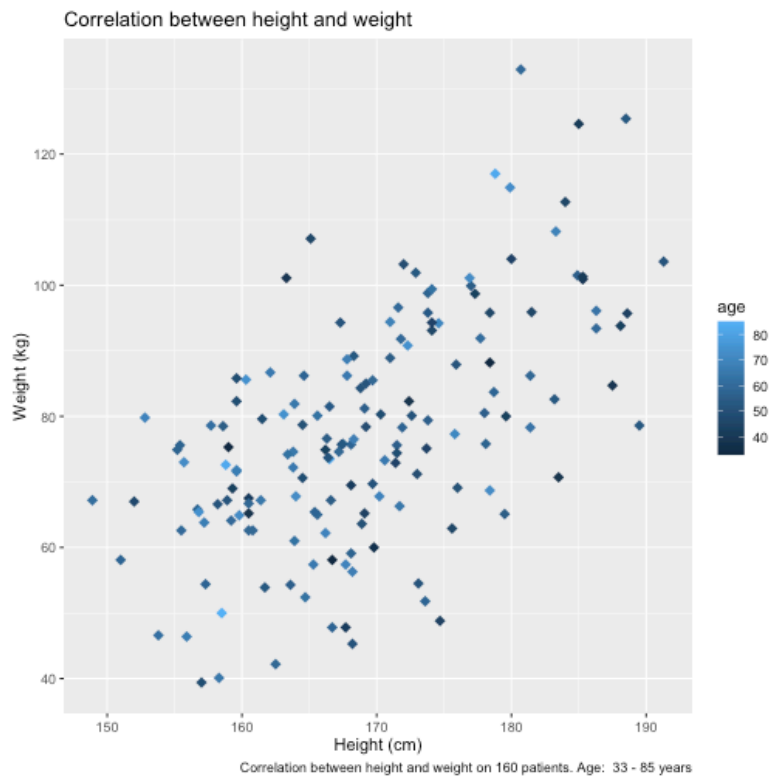

3.4 Scales

The definition of scales from the official documentation is "Scales control the details of how data values are translated to visual properties". Practically speaking, with scales we can control the data on the axis, the colors and the other parameters of the aesthetics (linetype, size, fill) by calling the correspondent function (that's what "" meant in the mock code).

```
ggplot() +
  geom_<type>() +
  labs() +
  scale_<type>()
```

We will see the main application of it.

3.4.1 Axis

With `scales_x/y_<type>` (if continuous, binned, manual or discrete) we can control the details of the axis, such as breaks, labels, limits, mathematical transformation and so on.

For example, we can change the breaks on both x and y axis of diastolic_bp_boxplot:

- on y axis, the distance between breaks is now 10, we change it to 5
- on x axis, we change the labels of the axis with the first letter uppercase

```
diastolic_bp_boxplot <- diastolic_bp_boxplot +
  labs(x = "Smoking", y = "Diastolic blood pressure (mmHg)") +
  scale_y_continuous(breaks = seq(0, max(df$diastolic_bp_mmhg), 5)) +
  scale_x_discrete(labels = c("Current", "Former", "Never"))
diastolic_bp_boxplot
```

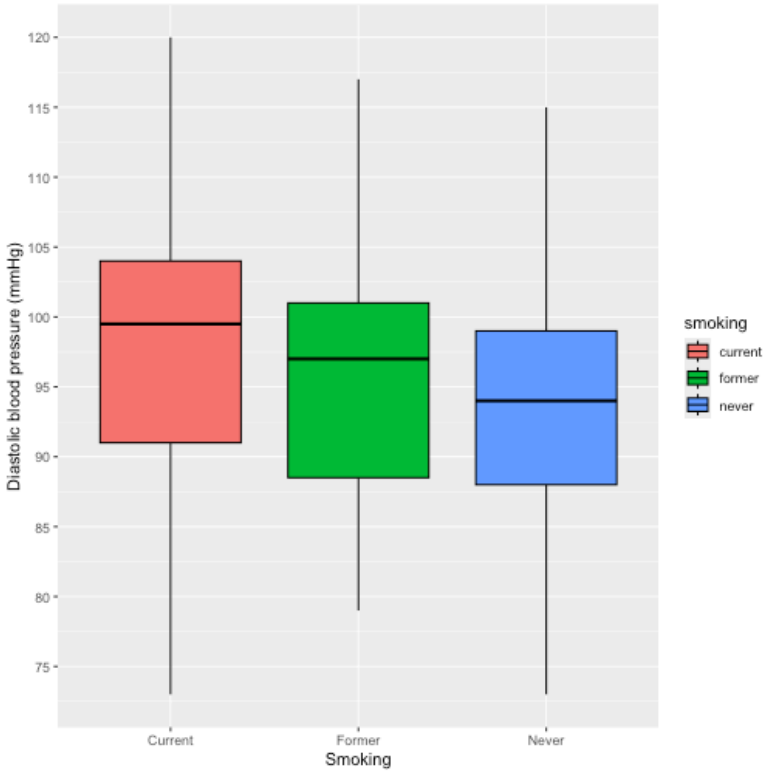

QUESTO SOLO NEL SITO

We can also apply mathematical transformation to continuous axis (like log scales):

```
BMI_stages_density +
  scale_x_continuous(trans = "log10")
```

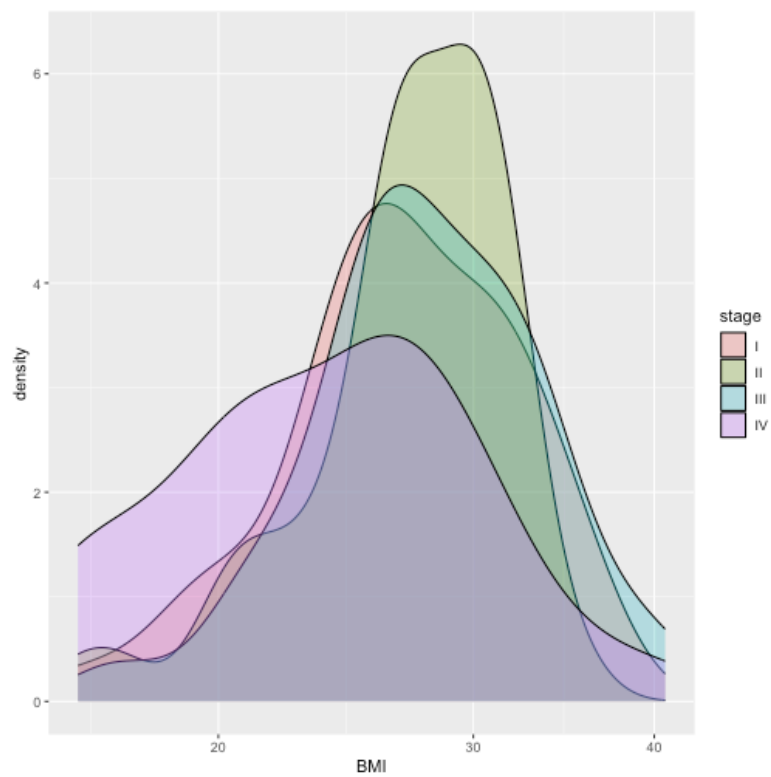


```
BMI_stages_density +
  scale_x_log10()
```

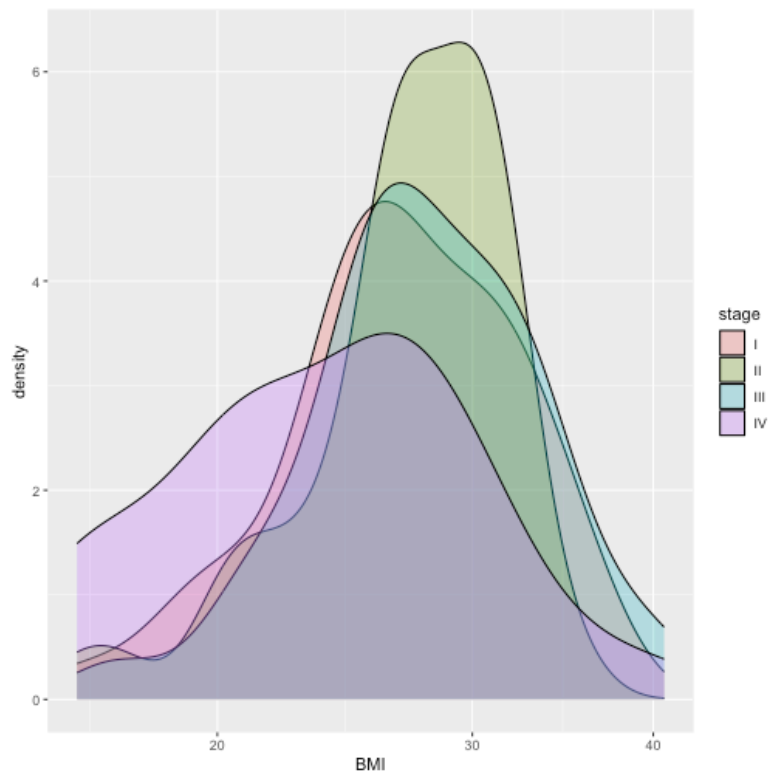

3.4.2 Colors

When we set `fill` or `color` aesthetics, we actually do not choose the colors that ggplot apply, we just say to it to make them different. With scales, we can set the colors we want. This site is very useful to choose colors.

For example, in stageI_female_diastolic_means_barplot, we want the Control to be gray, Drug_A to be blueish and Drug_B to be red. We can set it up like this:

```
stageI_female_diastolic_means_barplot <- stageI_female_diastolic_means_barplot +
  scale_fill_manual(values = c("Control" = "#568A81",
                               "Drug_A" = "#9FC2CC",
                               "Drug_B" = "#F95738"))
stageI_female_diastolic_means_barplot
```

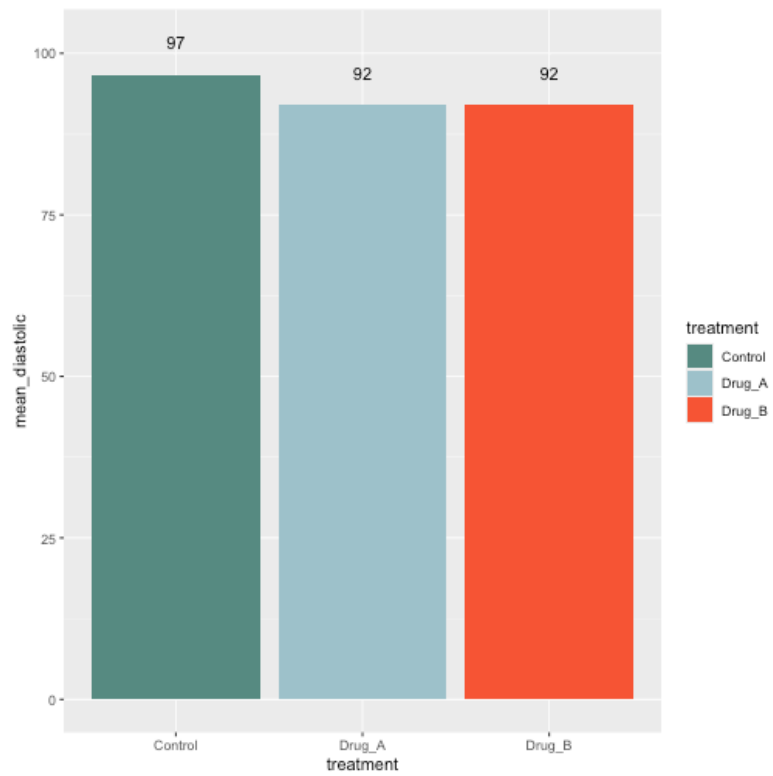

See this insite.) on hex color codes.

We can also change the color bar of corr_plot:

```
corr_plot <- corr_plot +
  scale_color_gradient(low = "#B8B7BA", high = "#3E2F5B")
corr_plot
```

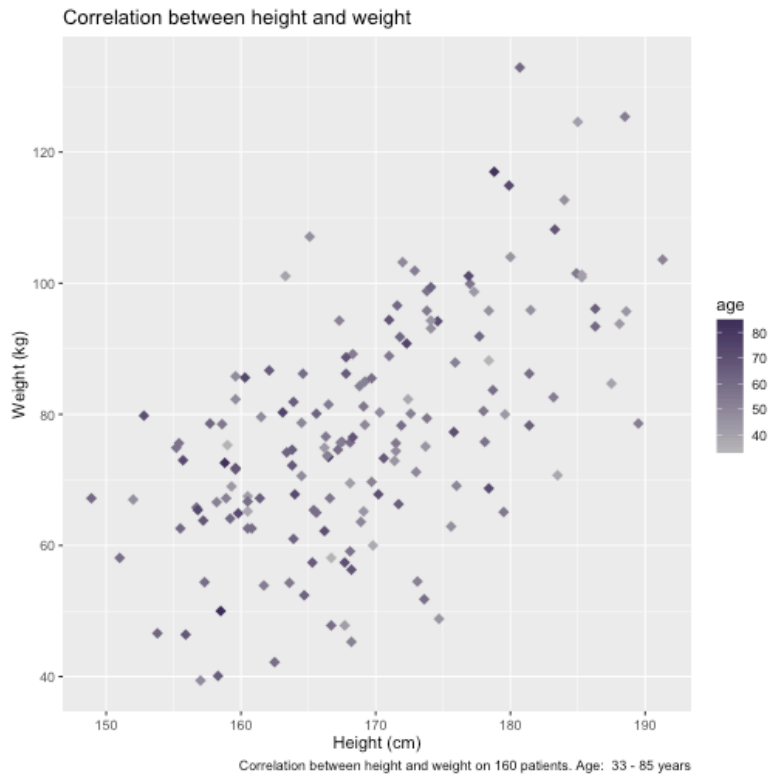

See this stackoverflow post for some useful tips and explanation on color scales in ggplot.

Exercise 3.1. Create a boxplot of platelets_10e9l in males, divided by stage (x axis) and treatment (fill). Then:

- y axis label should be "Platelets 10^9/L", with breaks every 75 starting from the minum value of platelets_10e9l up to the maximum value of platelets_10e9l
- Treatment colors should be the same as the ones in stageI_female_diastolic_means_barplot
- The title of the plot shoul be meaningful

Solution

```
males_df <- df |>
    filter(sex == "Male")

ggplot(males_df) +
    geom_boxplot(aes(x = stage, y = platelets_10e9l, fill = treatment)) +
    labs(y = "Platelets 10^9/L",
```

```
        title = "Distribution of Platelets concentration in males") +
    scale_fill_manual(values = c("Control" = "#568A81",
                                 "Drug_A" = "#9FC2CC",
                                 "Drug_B" = "#F95738")) +
    scale_y_continuous(breaks = seq(min(males_df$platelets_10e9l), max(males_df$platel
```

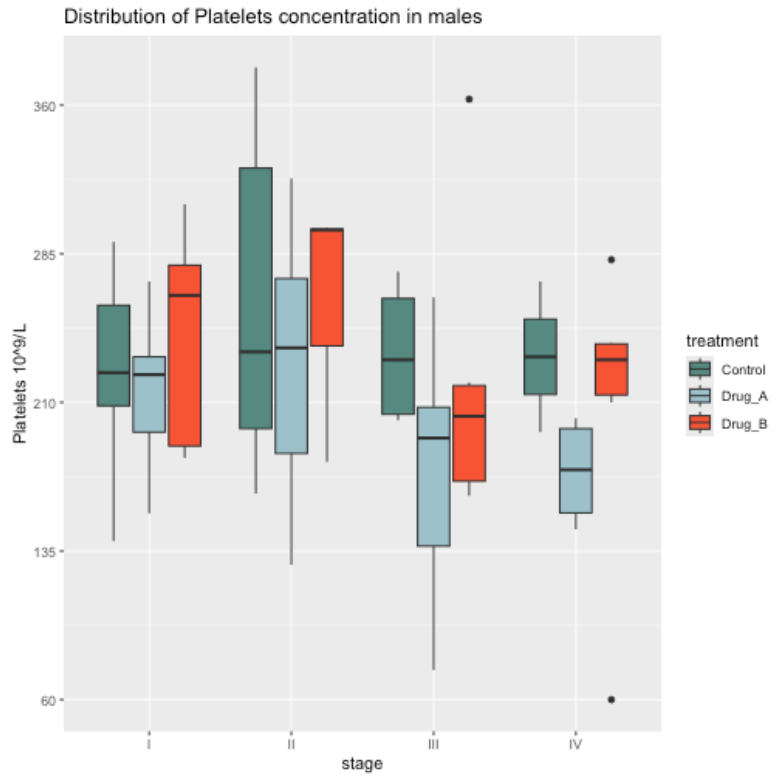

3.5 Facet

Faceting allow us to create a grid of panels, splitting data based on categorical variables.

```
ggplot() +
  geom_<type>() +
  labs() +
  scale_<type>() +
  facet_<type>() +
```

Let's see how we can rearrange the same data as in all_n_barplot but using multiple panels:

```
ggplot(df) +
  geom_bar(mapping = aes(x = stage, fill = treatment), position = "dodge") +
    facet_wrap(~ site)
```

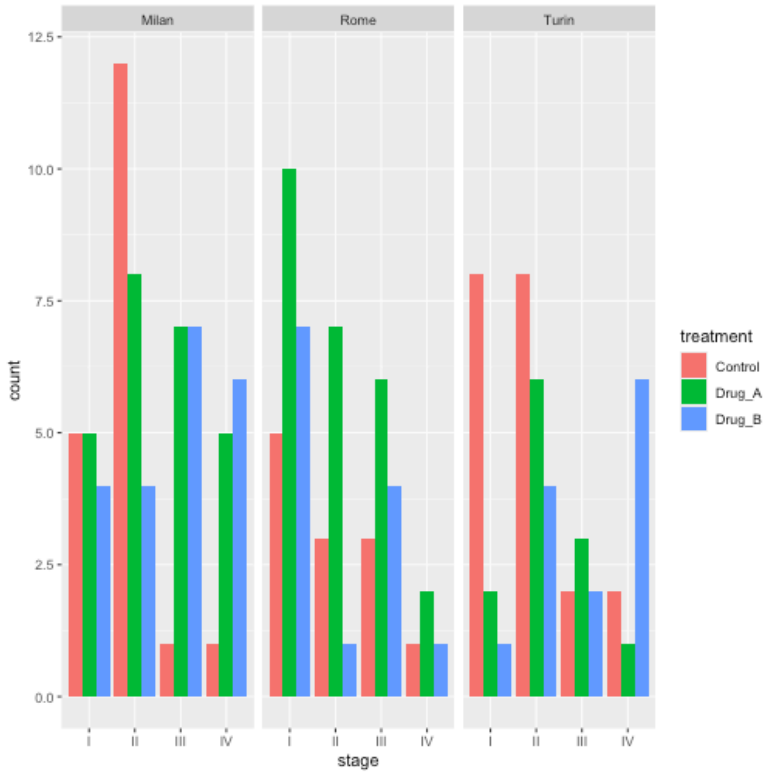

We can also customize the labels, similar to using scales:

```
corr_plot +
  facet_wrap(~ age > 55, labeller = labeller(`age > 55` = c("TRUE" = "> 55 yo", "FALSE" = "< 55 y
```

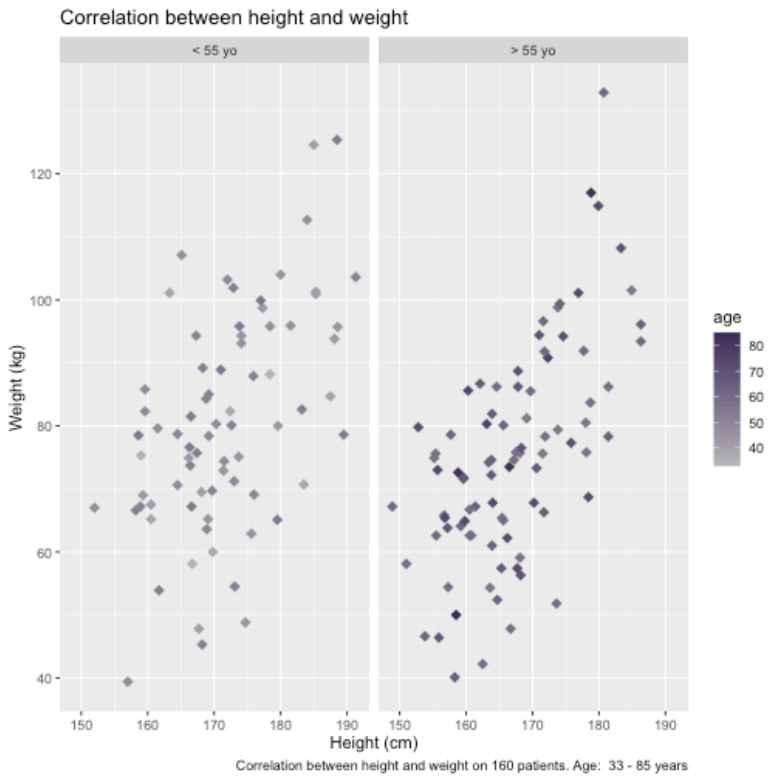

3.6 Theme

Ok, let's now see the last part (and maybe the most useful) of code. Theme allow us to control practically every single elements of the plot (except for the data): lines, font, labels, breaks, etc.

```
ggplot() +
  geom_<type>() +
  labs() +
  scale_<type>() +
  facet_<type>() +
  theme_<type>
```

There are few pre-installed theme, and we can further customize things with `theme`. Let's firt apply classic theme to diastolic_bp_boxplot:

```
diastolic_bp_boxplot <- diastolic_bp_boxplot +
  theme_classic()
diastolic_bp_boxplot
```

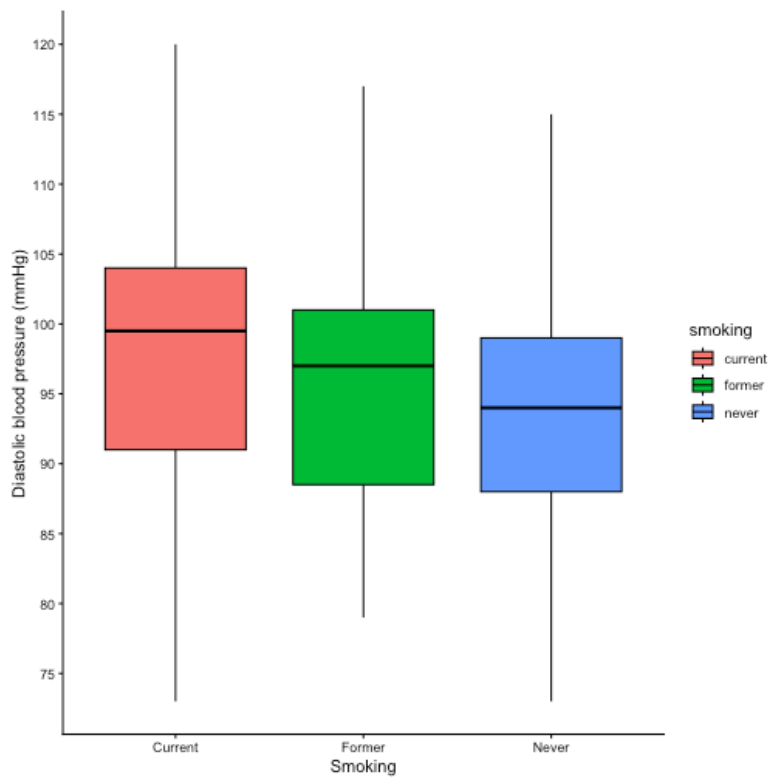

And now let's customize corr_plot:

- Title in bold and 16 pt size
- Add major grid lines in gray and 0.3 pt size
- Rotate x axis labels to 45°
- Change the adjustment of plot caption to left
- Add bottom margin to x title

```
corr_plot <- corr_plot +
  theme_classic() +
  theme(plot.title = element_text(face = "bold", size = 16),
        panel.grid.major = element_line(colour = "gray", linewidth = 0.3),
        axis.text.x = element_text(angle = 45, hjust = 1),
        plot.caption = element_text(hjust = 0),
```

```
        axis.title.x = element_text(margin = margin(b = 12)))
corr_plot
```

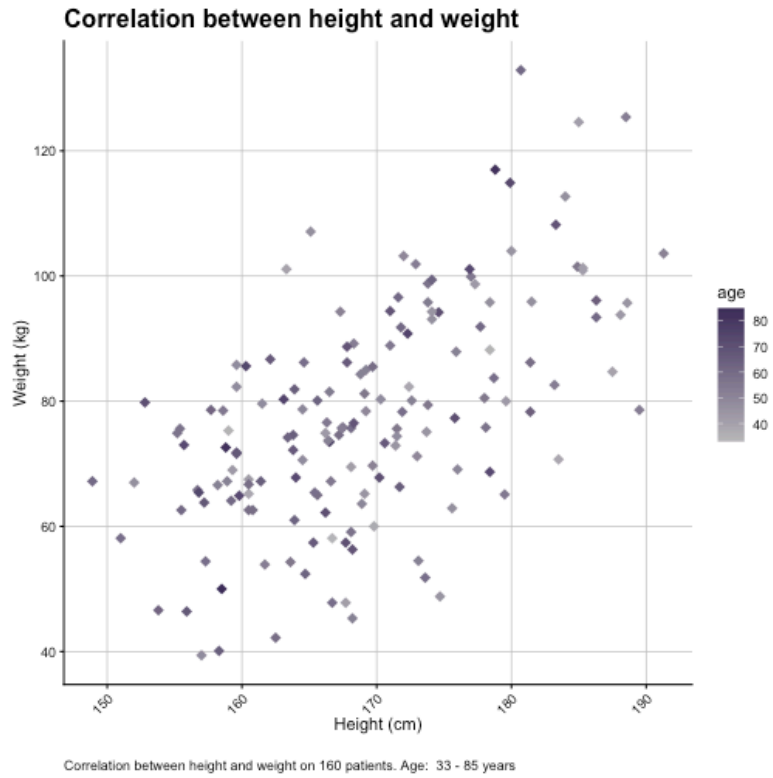

3.7 Save the plots

To save a plot to a file, we can use the function `ggsave` like this:

```
ggsave(plot = corr_plot, filename = "output/height_weight_corr-plot.pdf", units = "mm"
```

We can set units, width and height. See this article for further details about dimensions, dpi, dots etc.

3.8 Advanced examples

Here are some more advanced examples:

```
BMI_stages_density +
  labs(title = "BMI distribution at different stages", fill = "Stage") + # change labels
  scale_y_continuous(expand = expansion(mult = c(0, 0.1))) + # remove space at the bottom of y as
  scale_x_continuous(expand = expansion(mult = c(0, 0)), breaks = seq(0, max(df$BMI), 2.5)) + # r
  scale_fill_manual(values = c("I" = "#E0A890", "II" = "#2B59C3", "III" = "#9E2B25", "IV" = "#23C
  theme_classic() +
  theme(axis.line.y = element_blank(), # remove y-axes line
        axis.text.y = element_blank(), # remove y-axes labels
        axis.ticks.y = element_blank(), # remove y-axes ticks
        axis.title.y = element_blank(), # remove y-axes title
        legend.position = "top", # change legend position
        plot.title = element_text(size = 18, family = "sans", face = "bold", margin = margin(b =
        axis.title.x = element_text(size = 12, margin = margin(t = 10)),
        legend.key.size = unit(x = 5, units = "mm") # change legens key size
        )
```

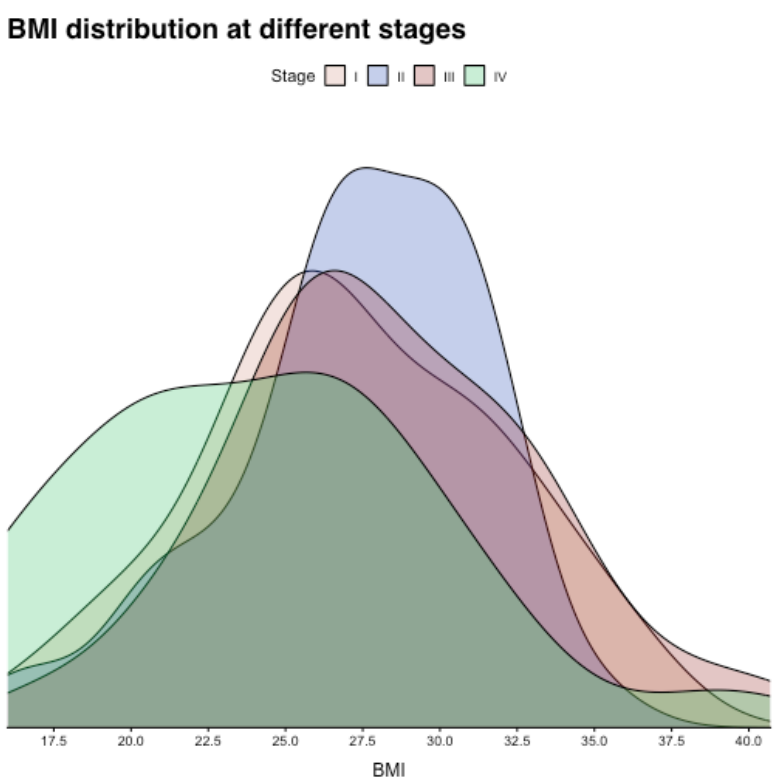


```
ggplot(df) +
  stat_summary(aes(x = stage, y = height_cm, color = weight_kg > 90), fun.data = "mean_se", geom
  labs(y = "Height (cm)", color = "Weight") +
```

```
  scale_color_manual(labels = c("TRUE" = "> 90kg", "FALSE" = "< 90kg"), values = c("TRU
  theme_classic()
```

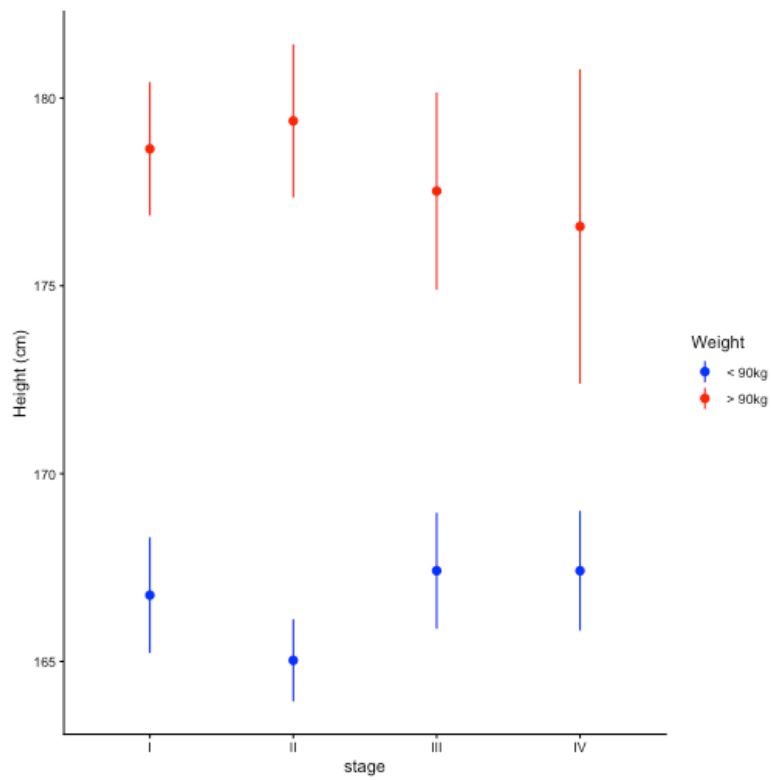

Exercise 3.2. Create a plot that shows if the levels of creatinine_mgdl changes depending on the year in each site. Set a proper scale for y axis, split the plot in panels, change the colors of the years, add a title. **Tip**: visit_year now is numeric, while we want it to be……..

bonus add the information about the numerosity in each group (year/site) at the bottom of x axis.

Solution

```
df$visit_year <- factor(df$visit_year)

df_n <- df |>
    group_by(site, visit_year) |>
    summarize(n = n(), .groups = "drop")

ggplot(df) +
```
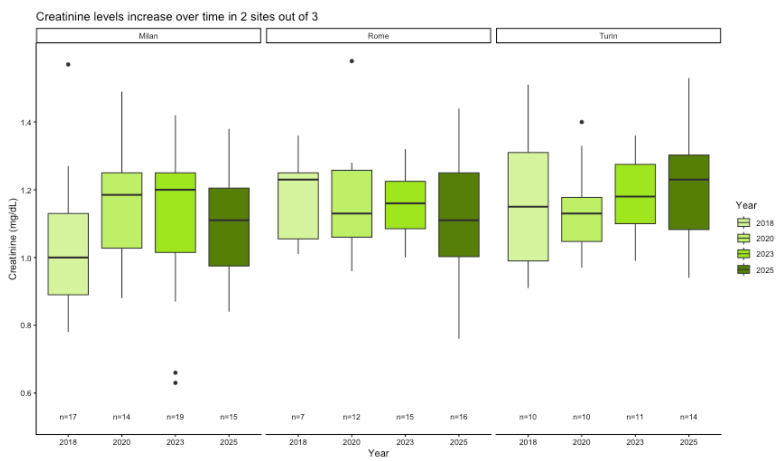
```
    geom_boxplot(aes(x = visit_year, y = creatinine_mgdl, fill = visit_year),
                 position = position_dodge(width = 0.75)) +
    geom_text(data = df_n,
              aes(x = visit_year, y = min(df$creatinine_mgdl) - 0.1, label = paste0("n=", n)),
              size = 3) +
    labs(y = "Creatinine (mg/dL)",
         title = "Creatinine levels increase over time in 2 sites out of 3",
         x = "Year",
         fill = "Year") +
    scale_fill_manual(values = c("2018" = "#D6F49D",
                                 "2020" = "#C1EF6C",
                                 "2023" = "#A2E723",
                                 "2025" = "#59810E")) +
    scale_y_continuous(breaks = seq(0, max(df$creatinine_mgdl), 0.2)) +
    facet_wrap(~ site) +
    theme_classic()
```